

# KingCountySales

August 23, 2022

## 0.1 King County House Regression - Project#2

- Student name: *Aysu Erdemir*
- Student pace: *Flex*
- Scheduled project review date/time: *August, 2022*
- Instructor name: *Claude Fried*

## 0.2 Overview

- This project analyzes house sale data from King County WA to provide insights and recommendations about the kind of houses “We Buy Ugly Houses” should invest on for their business.

## 0.3 Business Problem

- We Buy Ugly Houses is a real estate investor (House Flipper) that operates in King County WA. They purchase properties with the intention of remodeling to add value, then resell those properties for a profit.
- They want to know what type of houses to invest on for higher profit.

## 0.4 Data

King County House Sales dataset from [Kaggle](#) which contains house sale prices for King County sold between May 2014 and May 2015.

The variables/features included in the dataset are:

- **id** - Unique identifier for a house
- **date** - Date house was sold
- **price** - Sale price (prediction target)
- **bedrooms** - Number of bedrooms
- **bathrooms** - Number of bathrooms
- **sqft\_living** - Square footage of living space in the home
- **sqft\_lot** - Square footage of the lot
- **floors** - Number of floors (levels) in house
- **waterfront** - Whether the house is on a waterfront
- **view** - Quality of view from house
- **condition** - How good the overall condition of the house is. Related to maintenance of house.
  - See the [King County Assessor Website](#) for further explanation of each condition code
- **grade** - Overall grade of the house. Related to the construction and design of the house.

- See the [King County Assessor Website](#) for further explanation of each building grade code
- `sqft_above` - Square footage of house apart from basement
- `sqft_basement` - Square footage of the basement
- `yr_built` - Year when house was built
- `yr_renovated` - Year when house was renovated
- `zipcode` - ZIP Code used by the United States Postal Service
- `lat` - Latitude coordinate
- `long` - Longitude coordinate
- `sqft_living15` - The square footage of interior housing living space for the nearest 15 neighbors
- `sqft_lot15` - The square footage of the land lots of the nearest 15 neighbors

## 0.5 Methods

1. Clean the dataset.
2. Conduct feature engineering to come up with meaningful variables to be used in linear regression.
3. Build a series of linear regression models to come up with the best model to describe the relationship between the independent variables and the target/dependent variable (**house price**).
4. Check the linear regression assumptions to make sure normality, homoscedasticity are not violated and multicollinearity does not present.
5. Draw conclusions and make suggestions about the kind of houses to invest on.

```
[1]: # Import standard packages

import numpy as np
import pandas as pd
import statsmodels.api as sm
import scipy.stats as stats
from statsmodels.formula.api import ols
import statsmodels.formula.api as smf
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[2]: df = pd.read_csv("../data/kc_house_data.csv")
df.head()
```

```
[2]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
0	7129300520	10/13/2014	221900.0	3	1.00	1180	
1	6414100192	12/9/2014	538000.0	3	2.25	2570	
2	5631500400	2/25/2015	180000.0	2	1.00	770	
3	2487200875	12/9/2014	604000.0	4	3.00	1960	
4	1954400510	2/18/2015	510000.0	3	2.00	1680	

```
sqft_lot  floors  waterfront  view  ...      grade  sqft_above  \
```

0	5650	1.0	NaN	NONE	...	7 Average	1180
1	7242	2.0	NO	NONE	...	7 Average	2170
2	10000	1.0	NO	NONE	...	6 Low Average	770
3	5000	1.0	NO	NONE	...	7 Average	1050
4	8080	1.0	NO	NONE	...	8 Good	1680

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	\
0	0.0	1955	0.0	98178	47.5112	-122.257	
1	400.0	1951	1991.0	98125	47.7210	-122.319	
2	0.0	1933	NaN	98028	47.7379	-122.233	
3	910.0	1965	0.0	98136	47.5208	-122.393	
4	0.0	1987	0.0	98074	47.6168	-122.045	

	sqft_living15	sqft_lot15
0	1340	5650
1	1690	7639
2	2720	8062
3	1360	5000
4	1800	7503

[5 rows x 21 columns]

```
[3]: df.shape
```

```
[3]: (21597, 21)
```

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id              21597 non-null  int64
1   date            21597 non-null  object
2   price           21597 non-null  float64
3   bedrooms        21597 non-null  int64
4   bathrooms       21597 non-null  float64
5   sqft_living     21597 non-null  int64
6   sqft_lot        21597 non-null  int64
7   floors          21597 non-null  float64
8   waterfront      19221 non-null  object
9   view            21534 non-null  object
10  condition       21597 non-null  object
11  grade           21597 non-null  object
12  sqft_above      21597 non-null  int64
13  sqft_basement   21597 non-null  object
```

```

14 yr_built      21597 non-null  int64
15 yr_renovated  17755 non-null  float64
16 zipcode      21597 non-null  int64
17 lat          21597 non-null  float64
18 long         21597 non-null  float64
19 sqft_living15 21597 non-null  int64
20 sqft_lot15   21597 non-null  int64
dtypes: float64(6), int64(9), object(6)
memory usage: 3.5+ MB

```

```
[5]: df.isna().sum()
```

```

[5]: id          0
    date         0
    price        0
    bedrooms     0
    bathrooms    0
    sqft_living   0
    sqft_lot      0
    floors       0
    waterfront   2376
    view         63
    condition     0
    grade         0
    sqft_above    0
    sqft_basement 0
    yr_built      0
    yr_renovated  3842
    zipcode       0
    lat           0
    long          0
    sqft_living15 0
    sqft_lot15    0
    dtype: int64

```

- There are null values for waterfront, view and yr\_renovated.

```
[6]: df.duplicated().sum()
```

```
[6]: 0
```

```
[7]: df[df.duplicated('id')]
```

```

[7]:      id      date    price  bedrooms  bathrooms  sqft_living  \
94    6021501535  12/23/2014  700000.0         3         1.50         1580
314   4139480200  12/9/2014  1400000.0         4         3.25         4290
325   7520000520  3/11/2015  240500.0         2         1.00         1240
346   3969300030  12/29/2014  239900.0         4         1.00         1000

```

372	2231500030	3/24/2015	530000.0	4	2.25	2180
...	...	...	...	...	...	...
20165	7853400250	2/19/2015	645000.0	4	3.50	2910
20597	2724049222	12/1/2014	220000.0	2	2.50	1000
20654	8564860270	3/30/2015	502000.0	4	2.50	2680
20764	6300000226	5/4/2015	380000.0	4	1.00	1200
21565	7853420110	5/4/2015	625000.0	3	3.00	2780

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	\
94	5000	1.0	NO	NONE	...	8 Good	1290	
314	12103	1.0	NO	GOOD	...	11 Excellent	2690	
325	12092	1.0	NO	NONE	...	6 Low Average	960	
346	7134	1.0	NO	NONE	...	6 Low Average	1000	
372	10754	1.0	NO	NONE	...	7 Average	1100	
...	...	...	...	...	...	...	...	
20165	5260	2.0	NO	NONE	...	9 Better	2910	
20597	1092	2.0	NO	NONE	...	7 Average	990	
20654	5539	2.0	NaN	NONE	...	8 Good	2680	
20764	2171	1.5	NO	NONE	...	7 Average	1200	
21565	6000	2.0	NO	NONE	...	9 Better	2780	

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	\
94	290.0	1939	0.0	98117	47.6870	-122.386	
314	1600.0	1997	0.0	98006	47.5503	-122.102	
325	280.0	1922	1984.0	98146	47.4957	-122.352	
346	0.0	1943	NaN	98178	47.4897	-122.240	
372	1080.0	1954	0.0	98133	47.7711	-122.341	
...	...	...	...	...	...	...	
20165	0.0	2012	0.0	98065	47.5168	-121.883	
20597	10.0	2004	0.0	98118	47.5419	-122.271	
20654	0.0	2013	0.0	98045	47.4759	-121.734	
20764	0.0	1933	0.0	98133	47.7076	-122.342	
21565	0.0	2013	NaN	98065	47.5184	-121.886	

	sqft_living15	sqft_lot15
94	1570	4500
314	3860	11244
325	1820	7460
346	1020	7138
372	1810	6929
...	...	...
20165	2910	5260
20597	1330	1466
20654	2680	5992
20764	1130	1598
21565	2850	6000

[177 rows x 21 columns]

```
[8]: df[df['id'] == 6021501535]
```

```
[8]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
93	6021501535	7/25/2014	430000.0	3	1.5	1580	
94	6021501535	12/23/2014	700000.0	3	1.5	1580	

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	\
93	5000	1.0	NO	NONE	...	8 Good	1290	290.0	
94	5000	1.0	NO	NONE	...	8 Good	1290	290.0	

	yr_built	yr_renovated	zipcode	lat	long	sqft_living15	sqft_lot15
93	1939	0.0	98117	47.687	-122.386	1570	4500
94	1939	0.0	98117	47.687	-122.386	1570	4500

[2 rows x 21 columns]

- The same house was probably sold multiple times in the same year.
- Let's take only the most recent sell for those 177 duplicated house IDs.

## 0.6 DATA CLEANING:

### Drop Duplicates:

```
[9]: df = df.drop_duplicates(subset = 'id', keep = 'last').reset_index(drop=True)
df[df['id'] == 6021501535]
```

```
[9]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
93	6021501535	12/23/2014	700000.0	3	1.5	1580	

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	\
93	5000	1.0	NO	NONE	...	8 Good	1290	290.0	

	yr_built	yr_renovated	zipcode	lat	long	sqft_living15	sqft_lot15
93	1939	0.0	98117	47.687	-122.386	1570	4500

[1 rows x 21 columns]

```
[10]: df.shape
```

```
[10]: (21420, 21)
```

### Handling NaN values:

```
[11]: nulls = ['waterfront', 'view', 'yr_renovated']
print(*(f"{item}: {df[item].isnull().sum()}" for item in nulls), sep='\n' )
```

```
waterfront: 2353
view: 63
yr_renovated: 3813
```

```
[12]: df.waterfront.value_counts()
```

```
[12]: NO      18921
      YES      146
      Name: waterfront, dtype: int64
```

```
[13]: df.view.value_counts()
```

```
[13]: NONE      19253
      AVERAGE    956
      GOOD       505
      FAIR       329
      EXCELLENT   314
      Name: view, dtype: int64
```

```
[14]: df.yr_renovated.value_counts()
```

```
[14]: 0.0      16867
      2014.0     73
      2003.0     31
      2013.0     31
      2007.0     30
      ...
      1934.0      1
      1971.0      1
      1954.0      1
      1950.0      1
      1944.0      1
      Name: yr_renovated, Length: 70, dtype: int64
```

```
[15]: df.waterfront.isna().sum()/len(df)
```

```
[15]: 0.10985060690943044
```

- 11% of `waterfront` is NaN.
- Let's convert that to 0, because if a house had waterfront, it would likely be known and marked as YES.

```
[16]: df.view.isna().sum()/len(df.view)
```

```
[16]: 0.0029411764705882353
```

- .003 % of `view` is NaN.
- Let's convert that to NONE, because if a house had view, it would likely be known.

```
[17]: # We would expect houses with NaN on view also to be NaN or NO on waterfront,
      ↪let's double check that:
```

```
len(df[((df["waterfront"] == 'NO') | (df["waterfront"].isnull())) &
      (df["view"].isnull())])
# 62 out of 63 fits the criteria
```

```
[17]: 62
```

```
[18]: # Let's replace null with NONE
df['view'] = df['view'].fillna('NONE')
```

```
[19]: # We would expect houses with NaN on waterview also to be NONE on view, let's
      ↪double check that:
```

```
len(df[(df["view"] == 'NONE') & (df["waterfront"].isnull())])
# Majority of the data (2093/2353) fits the criteria.
```

```
[19]: 2093
```

```
[20]: # Let's replace null with NO
df['waterfront'] = df['waterfront'].fillna('NO')
```

```
[21]: df['yr_renovated'].describe()
      # Minimum is 0 most likely meaning that house has not been renovated.
      # We will replace all NaNs with 0 to mean the same thing.
```

```
[21]: count    17607.000000
      mean      83.890101
      std      400.534473
      min        0.000000
      25%        0.000000
      50%        0.000000
      75%        0.000000
      max      2015.000000
      Name: yr_renovated, dtype: float64
```

```
[22]: df['yr_renovated'] = df['yr_renovated'].fillna(0)
```

```
[23]: df.isnull().sum()
```

```
[23]: id          0
      date        0
      price       0
      bedrooms    0
      bathrooms   0
      sqft_living  0
      sqft_lot    0
```



```

floors          0
waterfront      0
view            0
condition       0
grade           0
sqft_above      0
sqft_basement   0
yr_built        0
yr_renovated    0
zipcode         0
lat             0
long            0
sqft_living15   0
sqft_lot15      0
dtype: int64

```

### Fixing variable types:

```
[24]: df_fixed = df.copy()
```

```
[25]: df_fixed.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21420 entries, 0 to 21419
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    21420 non-null  int64
1   date                  21420 non-null  object
2   price                 21420 non-null  float64
3   bedrooms              21420 non-null  int64
4   bathrooms             21420 non-null  float64
5   sqft_living           21420 non-null  int64
6   sqft_lot              21420 non-null  int64
7   floors                21420 non-null  float64
8   waterfront            21420 non-null  object
9   view                  21420 non-null  object
10  condition             21420 non-null  object
11  grade                 21420 non-null  object
12  sqft_above            21420 non-null  int64
13  sqft_basement         21420 non-null  object
14  yr_built              21420 non-null  int64
15  yr_renovated          21420 non-null  float64
16  zipcode               21420 non-null  int64
17  lat                   21420 non-null  float64
18  long                  21420 non-null  float64
19  sqft_living15         21420 non-null  int64
20  sqft_lot15            21420 non-null  int64

```

```
dtypes: float64(6), int64(9), object(6)
memory usage: 3.4+ MB
```

```
[26]: df_fixed.columns.to_series().groupby(df_fixed.dtypes).groups
      # object: ['date', 'waterfront', 'view', 'condition', 'grade', 'sqft_basement']}]
```

```
[26]: {int64: ['id', 'bedrooms', 'sqft_living', 'sqft_lot', 'sqft_above', 'yr_built',
          'zipcode', 'sqft_living15', 'sqft_lot15'], float64: ['price', 'bathrooms',
          'floors', 'yr_renovated', 'lat', 'long'], object: ['date', 'waterfront', 'view',
          'condition', 'grade', 'sqft_basement']}
```

These variables were coded as **string** and they need to be fixed (converted to **numerical**) for linear regression: - object: date, waterfront, view, condition, grade, sqft\_basement

```
[27]: df_fixed.waterfront.value_counts()
```

```
[27]: NO      21274
      YES      146
      Name: waterfront, dtype: int64
```

```
[28]: dic = {"NO":0, "YES":1}
      df_fixed.replace({"waterfront": dic}, inplace=True)
      df_fixed["waterfront"].value_counts()
```

```
[28]: 0      21274
      1      146
      Name: waterfront, dtype: int64
```

```
[29]: df_fixed["waterfront"].dtype
```

```
[29]: dtype('int64')
```

```
[30]: df_fixed['view'].value_counts()
```

```
[30]: NONE      19316
      AVERAGE    956
      GOOD       505
      FAIR       329
      EXCELLENT  314
      Name: view, dtype: int64
```

```
[31]: dic = {"NONE":1, "FAIR":2, "AVERAGE":3, "GOOD":4, "EXCELLENT":5}
      df_fixed.replace({"view": dic}, inplace=True)
      df_fixed["view"].value_counts()
```

```
[31]: 1      19316
      3       956
      4       505
```

```
2      329
5      314
Name: view, dtype: int64
```

```
[32]: df['condition'].value_counts()
```

```
[32]: Average      13900
      Good        5643
      Very Good   1687
      Fair         162
      Poor         28
      Name: condition, dtype: int64
```

```
[33]: dic = {"Poor":1, "Fair":2, "Average":3, "Good":4, "Very Good":5}
      df_fixed.replace({"condition": dic}, inplace=True)
      df_fixed["condition"].value_counts()
```

```
[33]: 3      13900
      4      5643
      5      1687
      2       162
      1        28
      Name: condition, dtype: int64
```

```
[34]: df_fixed['grade'].value_counts()
```

```
[34]: 7 Average      8889
      8 Good        6041
      9 Better      2606
      6 Low Average  1995
      10 Very Good   1130
      11 Excellent    396
      5 Fair         234
      12 Luxury        88
      4 Low          27
      13 Mansion      13
      3 Poor          1
      Name: grade, dtype: int64
```

```
[35]: dic = {"3 Poor":3, "4 Low":4, "5 Fair":5, "6 Low Average":6, "7 Average":7, "8_
      ↪Good":8, \
      "9 Better":9, "10 Very Good":10, "11 Excellent":11, "12 Luxury":12, "13_
      ↪Mansion":13}
      df_fixed.replace({"grade": dic}, inplace=True)
      df_fixed["grade"].value_counts()
```

```
[35]: 7      8889
      8      6041
      9      2606
      6      1995
     10      1130
     11       396
      5       234
     12        88
      4        27
     13        13
      3         1
      Name: grade, dtype: int64
```

```
[36]: df_fixed['sqft_basement'].unique()
```

```
[36]: array(['0.0', '400.0', '910.0', '1530.0', '?', '730.0', '1700.0', '300.0',
        '970.0', '760.0', '720.0', '700.0', '820.0', '780.0', '790.0',
        '330.0', '1620.0', '360.0', '588.0', '1510.0', '410.0', '990.0',
        '600.0', '560.0', '550.0', '1000.0', '1600.0', '500.0', '1040.0',
        '880.0', '1010.0', '240.0', '265.0', '290.0', '800.0', '540.0',
        '710.0', '840.0', '380.0', '770.0', '480.0', '570.0', '1490.0',
        '620.0', '1250.0', '1270.0', '120.0', '650.0', '180.0', '1130.0',
        '450.0', '1640.0', '1460.0', '1020.0', '1030.0', '750.0', '640.0',
        '1070.0', '490.0', '1310.0', '630.0', '2000.0', '390.0', '430.0',
        '850.0', '210.0', '1430.0', '1950.0', '440.0', '220.0', '1160.0',
        '860.0', '580.0', '2060.0', '1820.0', '1180.0', '200.0', '1150.0',
        '1200.0', '680.0', '530.0', '1450.0', '1170.0', '1080.0', '960.0',
        '280.0', '870.0', '1100.0', '460.0', '1400.0', '660.0', '1220.0',
        '900.0', '420.0', '1580.0', '1380.0', '475.0', '690.0', '270.0',
        '350.0', '935.0', '1370.0', '980.0', '1470.0', '160.0', '950.0',
        '50.0', '740.0', '1780.0', '1900.0', '340.0', '470.0', '370.0',
        '140.0', '1760.0', '130.0', '520.0', '890.0', '1110.0', '150.0',
        '1720.0', '810.0', '190.0', '1290.0', '670.0', '1800.0', '1120.0',
        '1810.0', '60.0', '1050.0', '940.0', '310.0', '930.0', '1390.0',
        '610.0', '1830.0', '1300.0', '510.0', '1330.0', '1590.0', '920.0',
        '1320.0', '1420.0', '1240.0', '1960.0', '1560.0', '2020.0',
        '1190.0', '2110.0', '1280.0', '250.0', '2390.0', '1230.0', '170.0',
        '830.0', '1260.0', '1410.0', '1340.0', '590.0', '1500.0', '1140.0',
        '260.0', '100.0', '320.0', '1480.0', '1060.0', '1284.0', '1670.0',
        '1350.0', '2570.0', '1090.0', '110.0', '2500.0', '90.0', '1940.0',
        '1550.0', '2350.0', '2490.0', '1481.0', '1360.0', '1135.0',
        '1520.0', '1850.0', '1660.0', '2130.0', '2600.0', '1690.0',
        '243.0', '1210.0', '1024.0', '1798.0', '1610.0', '1440.0',
        '1570.0', '1650.0', '704.0', '1910.0', '1630.0', '2360.0',
        '1852.0', '2090.0', '2400.0', '1790.0', '2150.0', '230.0', '70.0',
        '1680.0', '2100.0', '3000.0', '1870.0', '1710.0', '2030.0',
        '875.0', '1540.0', '2850.0', '2170.0', '506.0', '906.0', '145.0',
```

```
'2040.0', '784.0', '1750.0', '374.0', '518.0', '2720.0', '2730.0',
'1840.0', '3480.0', '2160.0', '1920.0', '2330.0', '1860.0',
'2050.0', '4820.0', '1913.0', '80.0', '2010.0', '3260.0', '2200.0',
'415.0', '1730.0', '652.0', '2196.0', '1930.0', '515.0', '40.0',
'2080.0', '2580.0', '1548.0', '1740.0', '235.0', '861.0', '1890.0',
'2220.0', '792.0', '2070.0', '4130.0', '2250.0', '2240.0',
'1990.0', '768.0', '2550.0', '435.0', '1008.0', '2300.0', '2610.0',
'666.0', '3500.0', '172.0', '1816.0', '2190.0', '1245.0', '1525.0',
'1880.0', '862.0', '946.0', '1281.0', '414.0', '2180.0', '276.0',
'1248.0', '602.0', '516.0', '176.0', '225.0', '1275.0', '266.0',
'283.0', '65.0', '2310.0', '10.0', '1770.0', '2120.0', '295.0',
'207.0', '915.0', '556.0', '417.0', '143.0', '508.0', '2810.0',
'20.0', '274.0', '248.0'], dtype=object)
```

```
[37]: # Replace ? with 0 and then convert to numerical data
df_fixed['sqft_basement'].replace('?', '0.0', inplace = True)
```

```
[38]: df_fixed['sqft_basement'] = pd.to_numeric(df_fixed['sqft_basement'])
df_fixed['sqft_basement'].dtype
```

```
[38]: dtype('float64')
```

```
[39]: df_fixed.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21420 entries, 0 to 21419
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     21420 non-null  int64
1   date                   21420 non-null  object
2   price                  21420 non-null  float64
3   bedrooms               21420 non-null  int64
4   bathrooms              21420 non-null  float64
5   sqft_living            21420 non-null  int64
6   sqft_lot               21420 non-null  int64
7   floors                 21420 non-null  float64
8   waterfront             21420 non-null  int64
9   view                   21420 non-null  int64
10  condition              21420 non-null  int64
11  grade                  21420 non-null  int64
12  sqft_above             21420 non-null  int64
13  sqft_basement          21420 non-null  float64
14  yr_built               21420 non-null  int64
15  yr_renovated           21420 non-null  float64
16  zipcode                21420 non-null  int64
17  lat                    21420 non-null  float64
```

```

18 long                21420 non-null float64
19 sqft_living15        21420 non-null int64
20 sqft_lot15           21420 non-null int64
dtypes: float64(7), int64(13), object(1)
memory usage: 3.4+ MB

```

- We are **not** changing `date` datatype because we will derive another variable from it, and then we will drop it.

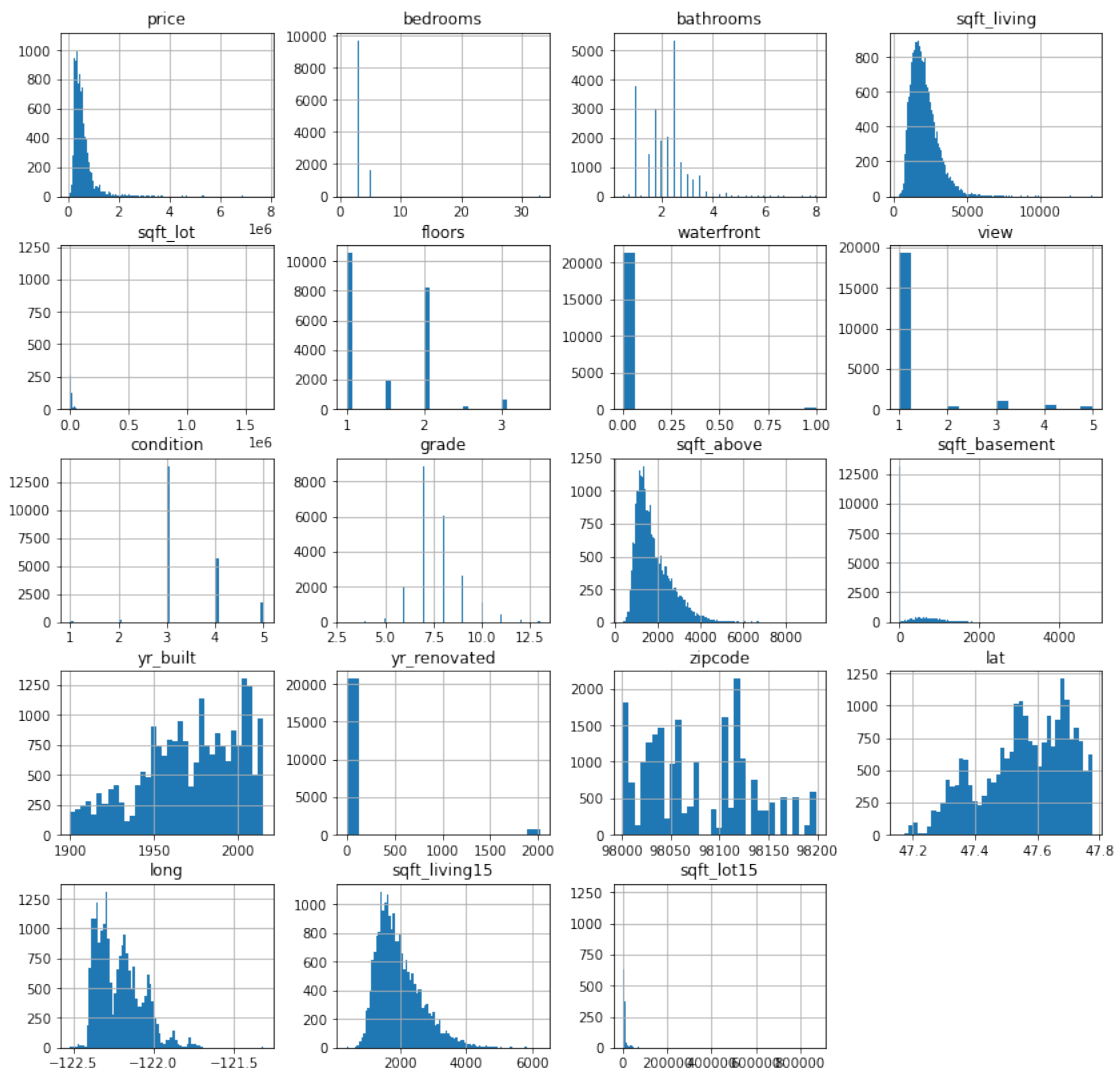
## 0.7 Feature Engineering:

```
[40]: df_new = df_fixed.copy()
```

- Drop `id` column since it has no meaning

```
[41]: df_new.drop(columns=['id'], inplace = True, axis=1)
```

```
[42]: df_new.hist(bins='auto', edgecolor='none', figsize=(14,14));
```



### The target / dependent variable:

- The sale prices have a negative skew, meaning the majority of the data is in the lower values, and there are fewer very high values.
- We might benefit from (log) transforming this variable.

```
[43]: # Code copied from: https://stackoverflow.com/questions/61330427/
      ↪ set-y-axis-in-millions and modified a bit.
      # Use the function below to get rid of 1e8s etc on graphs and to format numbers
      ↪ in thousands, millions, etc in visualizations...

from matplotlib.ticker import FuncFormatter

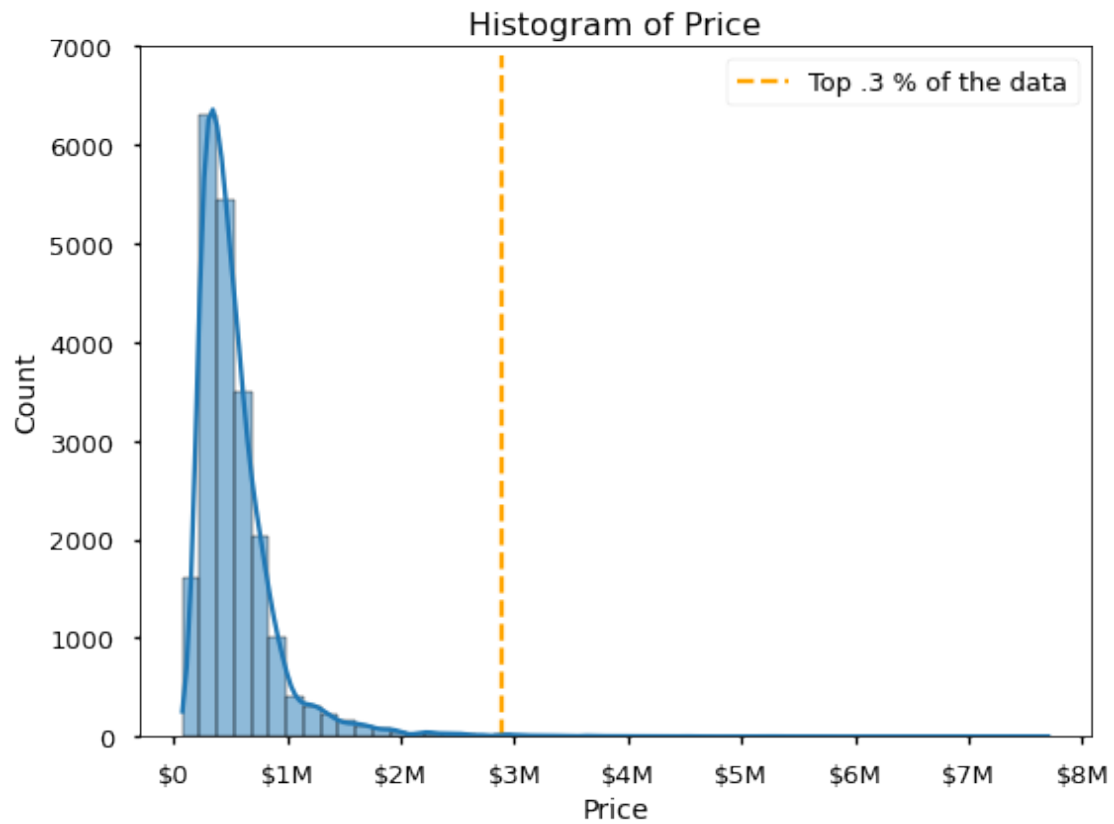
def human_format(num, pos):
    magnitude = 0
    while abs(num) >= 1000:
        magnitude += 1
        num /= 1000.0
    return '%%.0f%s' % (num, ['', 'K', 'M', 'B', 'T', 'P'][magnitude])

formatter = FuncFormatter(human_format)
```

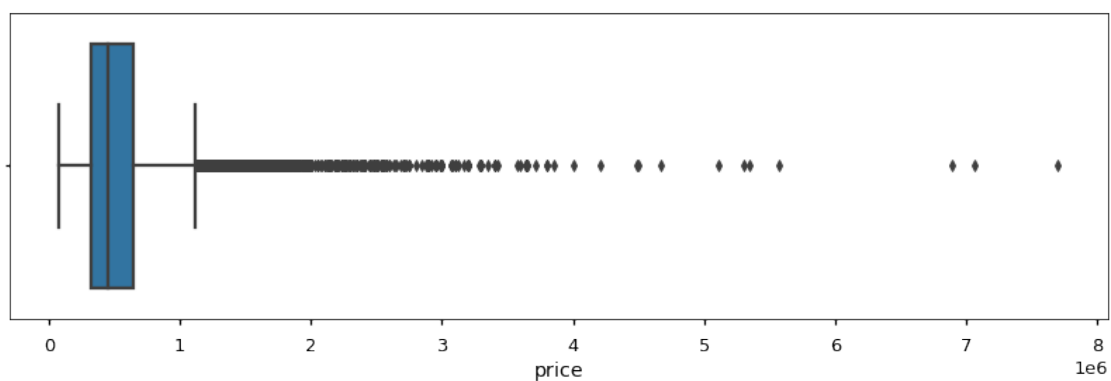
```
[44]: with plt.style.context('seaborn-talk'):
      fig, ax = plt.subplots(figsize=(8, 6))

      sns.histplot(x = df_new['price'], bins=50, ax=ax, kde = True)
      ax.xaxis.set_major_formatter(formatter)
      plt.ylim(0, 7000)
      ax.set_title('Histogram of Price', fontsize=16)
      ax.set_xlabel("Price", fontsize=14)
      ax.set_ylabel("Count", fontsize=14)
      ax.vlines(df['price'].quantile(0.997), 0, 6900, color= 'orange',
      ↪ linestyle='--', label = "Top .3 % of the data")
      ax.legend(loc = 'upper right')
      fig.tight_layout();

      fig.savefig('./images/Histogram_DependentMeasure.png', dpi=300)
```



```
[45]: with plt.style.context('seaborn-talk'):
      fig, ax = plt.subplots(figsize=(14, 4))
      sns.boxplot(x = df_new['price'], ax=ax);
```



**Remove the extreme outliers from the data:**

- Since the distribution is highly skewed I will remove the top .3% of the data.



- We will only be dealing with houses up to ~3M in price.


```
[46]: oldshape = df_new.shape
print(f"oldshape: {oldshape}")

df_new = df_new[(df_new['price'] < df_new['price'].quantile(.997))]
print(f"newshape: {df_new.shape}")
print(oldshape[0] - df_new.shape[0])
print(((oldshape[0] - df_new.shape[0]) * 100) / df_new.shape[0])

# we removed only 65 data points and .3 % of data.
```

```
oldshape: (21420, 20)
newshape: (21355, 20)
65
0.30437836572231325
```

```
[47]: with plt.style.context('seaborn-talk'):
fig, ax = plt.subplots(figsize=(15, 5))
sns.boxplot(x = df_new['price'], ax=ax)
```



### Log Transform the target variable:

- Price distribution is still highly skewed so let's log transform the variable as well.

```
[48]: df_new['log_price'] = np.log(df_new['price'])
```

```
[49]: with plt.style.context('seaborn-talk'):


fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(12, 6))
fig.set_tight_layout(True)

sns.histplot(x = df_new['price'], ax= ax1, bins=20);
```

```
sns.histplot(x = df_new['log_price'], ax= ax2, bins=20);

ax1.set_xlabel("Price", fontsize=14)
ax2.set_xlabel("Log Price", fontsize=14)
ax1.set_ylabel("Count", fontsize=14)
ax2.set_ylabel("Count", fontsize=14)

fig.savefig('./images/price_before_after_log.png', dpi=300);
```



- The price distribution looks NORMAL after log transformation.

### Creating a Binary View variable:

- The majority of the houses don't have a view and very few have ratings of 2-5.
- So let's impute a new variable to indicate whether or not a house has a view.

```
[50]: df_new['view'].value_counts()
```

```
[50]: 1    19297
      3     950
      4     499
      2     325
      5     284
      Name: view, dtype: int64
```

```
[51]: dic = {1:0, 2:1, 3:1, 4:1, 5:1}
      df_new['has_view'] = df_new['view'].map(dic)
      df_new['has_view'].value_counts()
```

```
[51]: 0    19297
      1     2058
      Name: has_view, dtype: int64
```

```
[52]: print(df_new.corr()['price']['view'])
      print(df_new.corr()['price']['has_view'])

      # Let's use `has_view` instead of `view`
```

```
0.37483613089845086
```


```
0.3500682990327732
```

```
[55]: with plt.style.context('seaborn-talk'):
      base_color = sns.color_palette("husl", 9)[6]
      fig, (ax1, ax2) = plt.subplots(ncols = 2, nrows=1, figsize=(12, 5))
      fig.set_tight_layout(True)

      sns.boxplot(x="waterfront", y="price", ax=ax1, data=df_new, color =_
↳base_color)
      ax1.yaxis.set_major_formatter(formatter)
      ax1.set_xticklabels(labels=['NO', 'YES'])
      ax1.set_title('Relationship Between Being Waterfront and Price',_
↳fontsize=14)
      ax1.set_xlabel("Waterfront",fontsize=14)
      ax1.set_ylabel("Price",fontsize=14)

      sns.boxplot(x="has_view", y="price", ax=ax2, data=df_new, color =_
↳base_color)
      ax2.yaxis.set_major_formatter(formatter)
      ax2.set_xticklabels(labels=['NO', 'YES'])
      ax2.set_title('Relationship Between Having a View and Price', fontsize=14)
      ax2.set_xlabel("View",fontsize=14)
      ax2.set_ylabel("Price",fontsize=14)

      fig.savefig('./images/view_waterfront _to_Price.png', dpi=300);
```



### Creating a Month variable:

- Date a house is sold cannot be significant in predicting house price so I will add a `month` column to indicate which month of the year the house was sold and drop the date.

```
[56]: df_new['month'] = pd.to_datetime(df_new['date']).dt.month
df_new.drop(columns=['date'], inplace = True, axis=1)
```

```
[57]: mean_month = pd.DataFrame(df_new.groupby('month')['price'].median()) # median
      ↳because price is skewed
mean_month.head()
```

```
[57]:      price
month
1      440000.0
2      426045.0
3      450000.0
4      475000.0
5      462000.0
```

```
[58]: mean_month = pd.DataFrame(df_new.groupby('month')['price'].median()) # median
      ↳because price is skewed


with plt.style.context('seaborn-talk'):
    base_color = [sns.color_palette("husl", 9)[0] if month == 4 else sns.
↳color_palette("husl", 9)[6] for month in mean_month.index]
    fig, ax = plt.subplots(figsize=(8, 6))
    #sns.barplot(x = mean_month.index, y = mean_month['price'], ax=ax, color =
↳base_color)
    bars = plt.bar(x=mean_month.index, height=mean_month['price'], color =
↳base_color)
```

```

plt.xticks(np.arange(1, 13, 1))
ax.set_xticklabels(labels=['January', 'February', 'March', 'April',
                           'May', 'June', 'July', 'August',
                           'September', 'October', 'November', 'December'],
rotation = 35)
ax.yaxis.set_major_formatter(formatter)
ax.set_title('Relationship Between Month of the Year House Sold and
Price',fontsize=14)
ax.set_xlabel("Month",fontsize=14)
ax.set_ylabel("Median Price",fontsize=14)
fig.tight_layout();

fig.savefig('./images/month_price-relationship.png', dpi=300);

```



```

[59]: dic = {1: 'january', 2: 'february', 3: 'march', 4: 'april', 5: 'may',
            6: 'june', 7: 'july', 8: 'august', 9: 'september',
            10: 'october', 11: 'november', 12: 'december'}

df_new.replace({"month": dic}, inplace=True)

```

### Dummy coding month variable:

- Months appear as distinct categories with no meaningful numerical relationship to one another.

```
[60]: month_dummies = pd.get_dummies(df_new['month']).drop(['january'], axis=1)
df_new = pd.concat([df_new, month_dummies], axis=1)
df_new = df_new.drop(['month'], axis=1)
df_new.head()
```

```
[60]:      price  bedrooms  bathrooms  sqft_living  sqft_lot  floors  waterfront  \
0  221900.0         3         1.00         1180      5650      1.0          0
1  538000.0         3         2.25         2570      7242      2.0          0
2  180000.0         2         1.00          770     10000      1.0          0
3  604000.0         4         3.00         1960      5000      1.0          0
4  510000.0         3         2.00         1680      8080      1.0          0

      view  condition  grade  ...  august  december  february  july  june  march  \
0       1          3       7  ...       0          0          0     0     0     0
1       1          3       7  ...       0          1          0     0     0     0
2       1          3       6  ...       0          0          1     0     0     0
3       1          5       7  ...       0          1          0     0     0     0
4       1          3       8  ...       0          0          1     0     0     0

      may  november  october  september
0       0          0         1          0
1       0          0         0          0
2       0          0         0          0
3       0          0         0          0
4       0          0         0          0
```

[5 rows x 32 columns]

### Creating an age related variable:

- Let's create a new variable called `age` to represent the age of an house from the time it was built or renovated using `yr_built` and `yr_renovated`.

```
[61]: # Because all houses were sold in 2014 and 2015 we will take 2015 as the
      current year.

df_new['age'] = 2015 - df_new['yr_built'] # Set all age based on yr_built
      initially.
mask = df_new['yr_renovated'] != 0 # create a mask for those rows with a value
      in yr_renovated.
df_new.loc[mask, "age"] = (2015 - df_new['yr_renovated']) # Set age based on
      yr_renovation where the mask condition is true
```

```
[62]: fig, ax = plt.subplots(figsize=(8, 6))
sns.regplot(x="age", y="price", ax=ax, data=df_new, color = base_color,
↳line_kws={"color": "orange"})
print(df_new.corr()['price']['age'])

# not much of a correlation but let's still keep this variable.
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-62-cb84da805dc2> in <module>
      1 fig, ax = plt.subplots(figsize=(8, 6))
----> 2 sns.regplot(x="age", y="price", ax=ax, data=df_new, color = base_color,
↳line_kws={"color": "orange"})
      3 print(df_new.corr()['price']['age'])
      4
      5 # not much of a correlation but let's still keep this variable.

/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/seaborn/_decorators.py
↳in inner_f(*args, **kwargs)
     44         )
     45         kwargs.update({k: arg for k, arg in zip(sig.parameters, args)})
----> 46         return f(**kwargs)
     47     return inner_f
     48

/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/seaborn/regression.py
↳in regplot(x, y, data, x_estimator, x_bins, x_ci, scatter, fit_reg, ci,
↳n_boot, units, seed, order, logistic, lowess, robust, logx, x_partial,
↳y_partial, truncate, dropna, x_jitter, y_jitter, label, color, marker,
↳scatter_kws, line_kws, ax)
     833     scatter_kws["marker"] = marker
     834     line_kws = {} if line_kws is None else copy.copy(line_kws)
--> 835     plotter.plot(ax, scatter_kws, line_kws)
     836     return ax
     837

/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/seaborn/regression.py
↳in plot(self, ax, scatter_kws, line_kws)
     357
     358         # Ensure that color is hex to avoid matplotlib weirdness
--> 359         color = mpl.colors.rgb2hex(mpl.colors.colorConverter.
↳to_rgb(color))
     360
     361         # Let color in keyword arguments override overall plot color

/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/matplotlib/colors.py
↳in to_rgb(c)
     344 def to_rgb(c):
```

```

345     """Convert *c* to an RGB color, silently dropping the alpha channel
↪ """
--> 346     return to_rgba(c)[:3]
347
348


/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/matplotlib/colors.py
↪ in to_rgba(c, alpha)
187         rgba = None
188         if rgba is None: # Suppress exception chaining of cache lookup
↪ failure.
--> 189         rgba = _to_rgba_no_colorcycle(c, alpha)
190         try:
191             _colors_full_map.cache[c, alpha] = rgba

/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/matplotlib/colors.py
↪ in _to_rgba_no_colorcycle(c, alpha)
263         raise ValueError(f"Invalid RGBA argument: {orig_c!r}")
264         if len(c) not in [3, 4]:
--> 265             raise ValueError("RGBA sequence should have length 3 or 4")
266         if not all(isinstance(x, Number) for x in c):
267             # Checks that don't work: `map(float, ...)` , `np.array(..., float)`
↪ and

ValueError: RGBA sequence should have length 3 or 4

```





- Let's create a new binary variable `age<30` where we group older versus younger houses.
- We pick age 30 as the criterion for a house that most likely needs repair.

```
[63]: df_new['age<30'] = df_new['age'] < 30  
df_new['age<30'].value_counts()
```

```
[63]: False    12809  
      True     8546  
      Name: age<30, dtype: int64
```

```
[64]: dic = {False:"0", True:"1"}  
df_new.replace({"age<30": dic}, inplace=True)  
df_new["age<30"] = df_new["age<30"].astype(int)  
df_new["age<30"].value_counts()
```

```
[64]: 0    12809  
      1    8546  
      Name: age<30, dtype: int64
```

```
[65]: print(df_new.corr()['price']['age'])  
      print(df_new.corr()['price']['age<30'])
```

```
-0.09219731443044772  
0.16020160020637306
```

- Improvement in correlation with price. Let's pick `age<30` over `age`.

```
[66]: df_new = df_new.drop(['yr_built', 'yr_renovated', 'age', 'view'], axis=1)
```


### Creating a location based variable:

- There are 70 Zipcodes! Too many levels if we go for One Hot Encoding.
- We cannot leave it as label encoded either, zip numbers do not have a meaningful numerical relationship to one another.
- I will engineer a new “location” variable with more meaningful categorical distinctions using Longitude-Latitude - I will then dummy code this variable.

```
[67]: zipmeans = df_new.groupby('zipcode')['price'].mean().  
      ↪sort_values(ascending=False)  
zipmeans = pd.DataFrame(zipmeans).reset_index()  
zipmeans.head()
```

```
[67]:   zipcode      price  
0    98039  1.761750e+06  
1    98004  1.245386e+06  
2    98040  1.129398e+06  
3    98112  1.051826e+06  
4    98006  8.418961e+05
```

```
[68]: with plt.style.context('seaborn-talk'):  
      base_color = sns.color_palette("husl", 9)[6]  
      fig, ax = plt.subplots(figsize=(8, 14))  
      sns.barplot(x = zipmeans['price'], y = zipmeans['zipcode'],  
                  order = zipmeans.sort_values('price', ascending = False).  
                  ↪zipcode, ax=ax, color = base_color, orient = "h")
```



The most expensive 4 zipcodes: - Medina, Bellevue, Mercer Island and Seattle

### Upload city data to use on the map:

- Upload US Zip Codes Database [here](#) which contains city info in relation to zipcodes:

```
[69]: dfzip = pd.read_csv("./data/usziips.csv")
dfzip.head()
```

```
[69]:  zip      lat      lng      city state_id  state_name  zcta  \
0  601  18.18027 -66.75266  Adjuntas      PR  Puerto Rico  True
1  602  18.36075 -67.17541   Aguada      PR  Puerto Rico  True
2  603  18.45744 -67.12225 Aguadilla      PR  Puerto Rico  True
3  606  18.16585 -66.93716   Maricao      PR  Puerto Rico  True
4  610  18.29110 -67.12243   Anasco      PR  Puerto Rico  True

      parent_zcta  population  density  county_fips  county_name  \
0             NaN      16773.0     100.5        72001   Adjuntas
1             NaN      37083.0     472.1        72003    Aguada
2             NaN      45652.0     513.2        72005  Aguadilla
3             NaN       6231.0      54.3        72093    Maricao
4             NaN      26502.0     275.7        72011   Añasco

                                county_weights  \
0                                {"72001": 98.76, "72141": 1.24}
1                                {"72003": 100}
2                                {"72005": 99.76, "72099": 0.24}
3    {"72093": 82.28, "72153": 11.67, "72121": 6.05}
4    {"72011": 96.71, "72099": 2.81, "72083": 0.37,...

      county_names_all      county_fips_all  imprecise  \
0  Adjuntas|Utuado      72001|72141      False
1           Aguada      72003      False
2  Aguadilla|Moca      72005|72099      False
3  Maricao|Yauco|Sabana Grande      72093|72153|72121      False
4  Añasco|Moca|Las Marías|Aguada  72011|72099|72083|72003      False

      military      timezone
0     False  America/Puerto_Rico
1     False  America/Puerto_Rico
2     False  America/Puerto_Rico
3     False  America/Puerto_Rico
4     False  America/Puerto_Rico
```

```
[70]: # subsetting the dataset to include those cities in KingCounty only:
```

```
dfzip = dfzip[(dfzip['county_names_all'].str.contains('King')) &
↳(dfzip['state_id'] == 'WA') ]
print(dfzip.zip.nunique())
dfzip
```

89

```
[70]:
```

	zip	lat	lng	city	state_id	state_name	zcta	\
32938	98001	47.30919	-122.26426	Auburn	WA	Washington	True	
32939	98002	47.30820	-122.21567	Auburn	WA	Washington	True	
32940	98003	47.30596	-122.31465	Federal Way	WA	Washington	True	
32941	98004	47.61865	-122.20548	Bellevue	WA	Washington	True	
32942	98005	47.61494	-122.16814	Bellevue	WA	Washington	True	
...	...	...	...	...	...	...	...	
33031	98199	47.65139	-122.40223	Seattle	WA	Washington	True	
33041	98224	47.73570	-121.56859	Baring	WA	Washington	True	
33092	98288	47.65204	-121.35740	Skykomish	WA	Washington	True	
33132	98354	47.25113	-122.31557	Milton	WA	Washington	True	
33178	98422	47.28907	-122.39123	Tacoma	WA	Washington	True	

	parent_zcta	population	density	county_fips	county_name	\
32938	NaN	34455.0	713.9	53033	King	
32939	NaN	33947.0	1829.6	53033	King	
32940	NaN	49445.0	1659.9	53033	King	
32941	NaN	37265.0	1979.1	53033	King	
32942	NaN	21414.0	1126.7	53033	King	
...	...	...	...	...	...	
33031	NaN	23444.0	2137.3	53033	King	
33041	NaN	243.0	1.5	53033	King	
33092	NaN	225.0	0.3	53033	King	
33132	NaN	7551.0	1029.0	53053	Pierce	
33178	NaN	21732.0	1197.8	53053	Pierce	

	county_weights	county_names_all	county_fips_all	\
32938	{"53033": 100}	King	53033	
32939	{"53033": 100}	King	53033	
32940	{"53033": 100}	King	53033	
32941	{"53033": 100}	King	53033	
32942	{"53033": 100}	King	53033	
...	...	...	...	
33031	{"53033": 100}	King	53033	
33041	{"53033": 100}	King	53033	
33092	{"53033": 100}	King	53033	
33132	{"53053": 80.02, "53033": 19.98}	Pierce King	53053 53033	
33178	{"53053": 97.78, "53033": 2.22}	Pierce King	53053 53033	

	imprecise	military	timezone
32938	False	False	America/Los_Angeles
32939	False	False	America/Los_Angeles
32940	False	False	America/Los_Angeles
32941	False	False	America/Los_Angeles
32942	False	False	America/Los_Angeles
...	...	...	...
33031	False	False	America/Los_Angeles
33041	False	False	America/Los_Angeles
33092	False	False	America/Los_Angeles
33132	False	False	America/Los_Angeles
33178	False	False	America/Los_Angeles

[89 rows x 18 columns]

```
[71]: # For cities with multiple zipcodes find an average location latitude and
      ↪ longitude
dfzip_table = dfzip.groupby('city')[['lat', 'lng']].mean()
dfzip_table = dfzip_table.reset_index()
dfzip_table
```

```
[71]:
```

	city	lat	lng
0	Auburn	47.301967	-122.203720
1	Baring	47.735700	-121.568590
2	Bellevue	47.602148	-122.155832
3	Black Diamond	47.311730	-122.003260
4	Bothell	47.757360	-122.198710
5	Carnation	47.696640	-121.840240
6	Duvall	47.740830	-121.934940
7	Enumclaw	47.171400	-121.679580
8	Fall City	47.573010	-121.902190
9	Federal Way	47.307385	-122.338315
10	Hobart	47.434410	-121.952400
11	Issaquah	47.530735	-122.005430
12	Kenmore	47.751620	-122.248920
13	Kent	47.382738	-122.191553
14	Kirkland	47.696140	-122.202995
15	Maple Valley	47.418610	-121.955890
16	Medina	47.633080	-122.239630
17	Mercer Island	47.566110	-122.232000
18	Milton	47.251130	-122.315570
19	North Bend	47.482760	-121.656780
20	Pacific	47.260790	-122.248220
21	Preston	47.547640	-121.936630
22	Ravensdale	47.339600	-121.890490
23	Redmond	47.673195	-122.070505

24	Renton	47.473778	-122.172132
25	Sammamish	47.604445	-122.041165
26	Seattle	47.593588	-122.329591
27	Skykomish	47.652040	-121.357400
28	Snoqualmie	47.592340	-121.789310
29	Snoqualmie Pass	47.451540	-121.357950
30	Tacoma	47.289070	-122.391230
31	Vashon	47.412190	-122.472600
32	Woodinville	47.757070	-122.094535

- We have a total of 32 cities in KingCounty WA which we can superimpose on the map below:

```
[72]: # For the purpose of graphing removing far east cities of Baring, Skykomish,
      ↪Snoqualmie Pass:
dfzip_table.drop(index=[1,27,29], axis=0, inplace=True)
```

```
[73]: # Scatterplot of longitude and latitude with a hue of price, city names are
      ↪superimposed to the map:
# Superimposed data is from: https://www.communitiescount.org/
      ↪king-county-geographies
# The shape is the shape of King County, WA


with plt.style.context('seaborn-talk'):

    fig, ax = plt.subplots(figsize=(12, 10))
    sns.scatterplot(data=dfzip_table, x='lng', y='lat', alpha = .5, ax=ax)
    [plt.text(x=row['lng'], y=row['lat'], s=row['city'], size='large',
      ↪color='black') for k,row in dfzip_table.iterrows()]

    sns.scatterplot(data=df_new, x='long', y='lat', hue='price',
      ↪size="price", sizes=(20, 200), alpha = .5, ax=ax)
    ax.axhline(y= 47.48, xmin=0, xmax=1, color='y', linestyle='--')
    ax.axvline(x= -122.25, ymin=0.52, ymax=0.9, color='y', linestyle='--')
    ax.axvline(x= -122.085, ymin=0.52, ymax=0.9, color='y', linestyle='--')
    ax.axhline(y= 47.745, xmin=0, xmax=1, color='y', linestyle='--')
    plt.xticks(np.arange(-122.6, -121.5, 0.1))

    plt.xlim(-122.6, -121.5)
    ax.set_title('Relationship Between Location and Price', fontsize=18)
    ax.set_xlabel("Longitude", fontsize=16)
    ax.set_ylabel("Latitude", fontsize=16)
    fig.tight_layout();

    fig.savefig('./images/LocationMap.png', dpi=300);
```



**Extract 5 regions based on coordinates:**

- North - Kenmore, Bothell, Woodinville
- East - Medina, Bellevue, Mercer Island and Kirkland
- Far East - Redmond, Sammamish etc.
- West - Seattle.
- South - Tacoma, Renton, Kent etc.

```
[74]: df_new['coordinates'] = list(zip(df_new.lat, df_new.long))

def region(coordinate):
    if (coordinate[0] > 47.745):
        return 'north'
    elif (coordinate[0] > 47.48) and (coordinate[0] < 47.745) and
    ↪(coordinate[1] < -122.25):
        return 'west'
    elif (coordinate[0] > 47.48) and (coordinate[0] < 47.745) and
    ↪(coordinate[1] > -122.25) and (coordinate[1] < -122.085):
```



```

        return 'east'
    elif (coordinate[0] > 47.48) and (coordinate[0] < 47.745) and
↪(coordinate[1] > -122.085):
        return 'fareast'
    else:
        return 'south'

region([47.5112, -122.257])

```

```
[74]: 'west'
```

```
[75]: df_new['region'] = df_new['coordinates'].apply(region)
df_new['region'].head()
```

```
[75]: 0      west
1      west
2      east
3      west
4  fareast
Name: region, dtype: object
```

```
[76]: df_new.groupby('region')['price'].median()
```

```
[76]: region
east      582250.0
fareast    585000.0
north     437000.0
south     299900.0
west      490000.0
Name: price, dtype: float64
```

```
[77]: df_new['region'].value_counts()
```

```
[77]: west      7344
south     5661
east      4448
fareast    2665
north     1237
Name: region, dtype: int64
```

```
[78]: mean_region = pd.DataFrame(df_new.groupby('region')['price'].median())
mean_region['price']
```


```
[78]: region
east      582250.0
fareast    585000.0
north     437000.0
```

```
south      299900.0
west       490000.0
Name: price, dtype: float64
```

```
[79]: mean_region = pd.DataFrame(df_new.groupby('region')['price'].median()) #
      ↪median because price is skewed

with plt.style.context('seaborn-talk'):
    base_color = sns.color_palette("husl", 9)[6]
    fig, ax = plt.subplots(figsize=(8, 6))
    sns.barplot(x = mean_region.index, y = mean_region['price'], ax=ax, color =
    ↪base_color)
    ax.set_xticklabels(labels=['East', 'Fareast', 'North', 'South',
    ↪'West-Seattle'])
    ax.yaxis.set_major_formatter(formatter)
    ax.set_title('Relationship Between Region and House Price',fontsize=14)
    ax.set_xlabel("Region",fontsize=14)
    ax.set_ylabel("Median Price",fontsize=14)
    fig.tight_layout();

    fig.savefig('./images/region_price-relationship.png', dpi=300);
```



- East regions appear as the most expensive area, followed by Seattle and then north. South is the cheapest.

```
[80]: # South region is the reference point
region_dummies = pd.get_dummies(df_new['region']).drop(['south'],axis=1)
df_new = pd.concat([df_new, region_dummies], axis=1)
df_new.head()
```

```
[80]:      price  bedrooms  bathrooms  sqft_living  sqft_lot  floors  waterfront  \
0  221900.0         3         1.00         1180     5650     1.0           0
1  538000.0         3         2.25         2570     7242     2.0           0
2  180000.0         2         1.00          770    10000     1.0           0
3  604000.0         4         3.00         1960     5000     1.0           0
4  510000.0         3         2.00         1680     8080     1.0           0
```

```
      condition  grade  sqft_above  ...  november  october  september  age<30  \
0             3       7         1180  ...         0         1           0         0
1             3       7         2170  ...         0         0           0         1
2             3       6          770  ...         0         0           0         0
3             5       7         1050  ...         0         0           0         0
4             3       8         1680  ...         0         0           0         1
```


```
      coordinates  region  east  fareast  north  west
0  (47.5112, -122.257)  west    0         0      0     1
1  (47.721, -122.319)  west    0         0      0     1
2  (47.7379, -122.233)  east    1         0      0     0
3  (47.5208, -122.393)  west    0         0      0     1
4  (47.6168, -122.045)  fareast  0         1      0     0
```

[5 rows x 36 columns]

```
[81]: df_new = df_new.drop(['region','coordinates', 'zipcode'], axis=1)
# Zipcodes cannot be left as label encoded since the levels do not have a
↳numerical relationship.
```

## 0.8 Feature Engineering Continued:

```
[82]: df_new.hist(bins='auto', edgecolor='none', figsize=(16,16));
```




**Remove outliers from bedrooms and bathrooms:**

```
[83]: fig, (ax1, ax2) = plt.subplots(ncols=2, nrows=1, figsize=(12, 4))
fig.set_tight_layout(True)

sns.histplot(x = df_new['bedrooms'], ax= ax1, bins=50);
sns.histplot(x = df_new['bathrooms'], ax =ax2, bins=50);

# There are outliers in both bathrooms and bedrooms.
```



```
[84]: df_new['bedrooms'].describe()
```

```
[84]: count    21355.000000
      mean      3.370265
      std      0.922473
      min      1.000000
      25%      3.000000
      50%      3.000000
      75%      4.000000
      max      33.000000
      Name: bedrooms, dtype: float64
```

```
[85]: df_new['bathrooms'].describe()
```

```
[85]: count    21355.000000
      mean      2.111672
      std      0.757100
      min      0.500000
      25%      1.750000
      50%      2.250000
      75%      2.500000
      max      8.000000
      Name: bathrooms, dtype: float64
```


```
[86]: print(df_new['bedrooms'].quantile(.999))
      print(df_new['bathrooms'].quantile(.999))
```

```
8.0
5.25
```

```
[87]: # Let's remove some very high values visible in the histogram ~ top 1 percent.
      df_new = df_new[df_new['bedrooms'] <= 8]
      df_new = df_new[df_new['bathrooms'] <= 6]
```

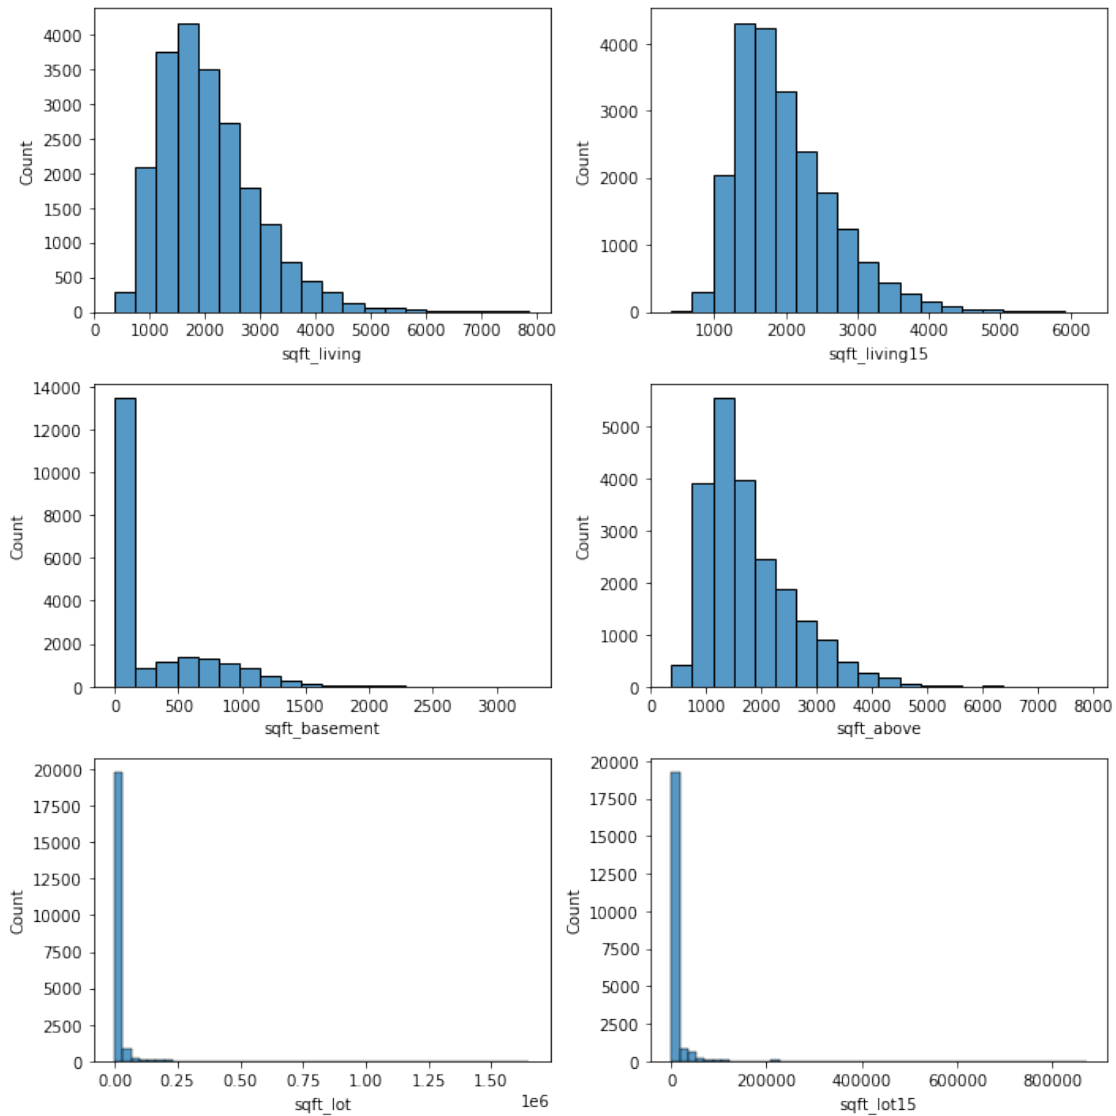
```
[88]: fig, (ax1, ax2) = plt.subplots(ncols=2, nrows=1, figsize=(12, 4))
fig.set_tight_layout(True)
```

```
sns.histplot(x = df_new['bedrooms'], ax= ax1, bins=10);
sns.histplot(x = df_new['bathrooms'], ax =ax2, bins=10);
```



```
[89]: fig, ((ax1, ax2),(ax3,ax4),(ax5,ax6)) = plt.subplots(ncols=2, nrows=3,
↳figsize=(10, 10))
fig.set_tight_layout(True)
```

```
sns.histplot(x = df_new['sqft_living'], ax= ax1, bins=20);
sns.histplot(x = df_new['sqft_living15'], ax= ax2, bins=20);
sns.histplot(x = df_new['sqft_basement'], ax =ax3, bins=20);
sns.histplot(x = df_new['sqft_above'], ax= ax4, bins=20);
sns.histplot(x = df_new['sqft_lot'], ax =ax5, bins=50);
sns.histplot(x = df_new['sqft_lot15'], ax= ax6, bins=50);
```



```
[90]: df_new.shape
```

```
[90]: (21340, 33)
```

```
[91]: print(df_new['sqft_lot'].quantile(.99)) # np.percentile(df_new['sqft_lot'],98)
      print(df_new['sqft_lot15'].quantile(.99))
```

```
213008.0
```


```
157687.0
```

Remove outliers from sqft\_lot, and sqft\_lot15:

```
[92]: # Remove some high values from sqft_lot and sqft_lot15
df_new = df_new[df_new['sqft_lot'] < 120000]
df_new = df_new[df_new['sqft_lot15'] < 120000]
```

```
[93]: fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(10, 4))
fig.set_tight_layout(True)

sns.histplot(x = df_new['sqft_lot'], ax= ax1, bins=20);
sns.histplot(x = df_new['sqft_lot15'], ax= ax2, bins=20);
```



Create a new `sqft_basement` variable:

- `has_basement` will define presence or absence of `sqft_basement` since more than half of the houses don't have a basement.

```
[94]: len(df_new[df_new['sqft_basement'] == 0])
```

```
[94]: 12834
```

```
[95]: df_new['has_basement'] = df_new['sqft_basement'] > 0
df_new['has_basement'].value_counts()
```

```
[95]: False    12834
      True     8070
      Name: has_basement, dtype: int64
```

```
[96]: dic = {False:"0", True:"1"}
df_new.replace({"has_basement": dic}, inplace=True)
df_new["has_basement"] = df_new["has_basement"].astype(int)
df_new["has_basement"].value_counts()
```



```
[96]: 0    12834
      1     8070
      Name: has_basement, dtype: int64
```

```
[97]: print(df_new.corr()['price']['sqft_basement'])
      print(df_new.corr()['price']['has_basement'])
      # correlation coef is smaller for has_basement but since this variable is more
      ↪ meaningful let's use it and drop 'sqft_basement'
```

```
0.2983845422940022
0.18230767995993896
```

```
[98]: len(df_new[df_new['sqft_lot'] == 0])
```

```
[98]: 0
```

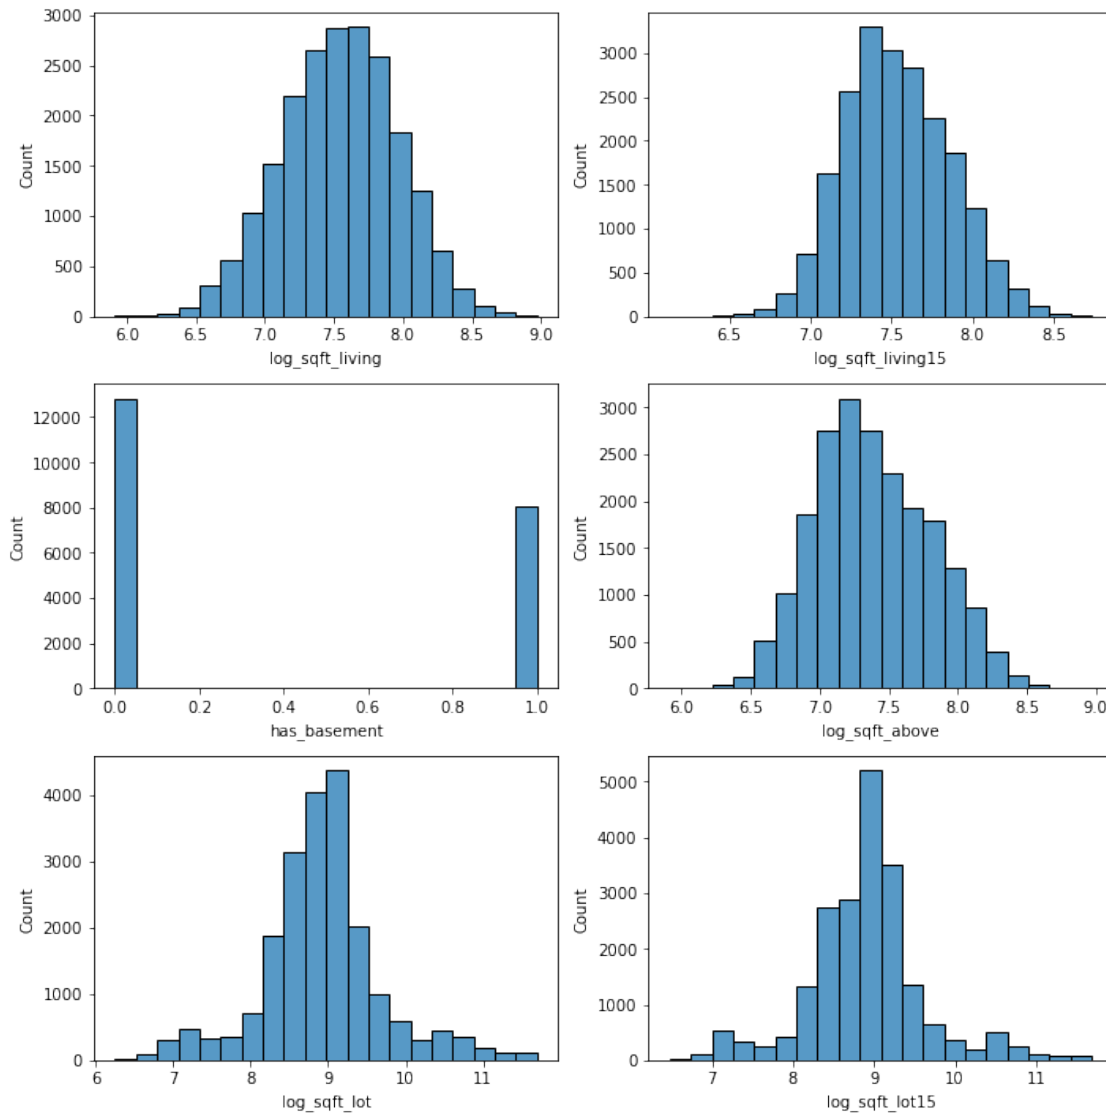
```
[99]: df_new.drop('sqft_basement', axis=1, inplace=True)
```

Log transform skewed variables in case we need to use them in regression:

```
[100]: for var in ['sqft_living', 'sqft_living15', 'sqft_above', 'sqft_lot', 'sqft_lot15']:
      df_new[f"log_{var}"] = np.log(df_new[var]) # df_new[f"log{var}"] = np.
      ↪ log(df_new[var])
```

```
[101]: fig, ((ax1, ax2), (ax3, ax4), (ax5, ax6)) = plt.subplots(ncols=2, nrows=3,
      ↪ figsize=(10, 10))
      fig.set_tight_layout(True)

      sns.histplot(x = df_new['log_sqft_living'], ax= ax1, bins=20);
      sns.histplot(x = df_new['log_sqft_living15'], ax= ax2, bins=20);
      sns.histplot(x = df_new['has_basement'], ax= ax3, bins=20);
      sns.histplot(x = df_new['log_sqft_above'], ax= ax4, bins=20);
      sns.histplot(x = df_new['log_sqft_lot'], ax =ax5, bins=20);
      sns.histplot(x = df_new['log_sqft_lot15'], ax= ax6, bins=20);
```



```
[102]: print(df.shape)
print(df_new.shape)
```

(21420, 21)

(20904, 38)

- In the end we lost  $\sim 2.5\%$  of the data during data engineering process:

```
[103]: total_data_loss = ((df_fixed.shape[0] - df_new.shape[0]) * 100) / df_fixed.
      ↪ shape[0]
total_data_loss
```

[103]: 2.4089635854341735

## 0.9 Feature Selection:

```
[104]: data = df_new.copy()
data.head()
```

```
[104]:      price  bedrooms  bathrooms  sqft_living  sqft_lot  floors  waterfront  \
0  221900.0          3         1.00         1180     5650     1.0           0
1  538000.0          3         2.25         2570     7242     2.0           0
2  180000.0          2         1.00          770    10000     1.0           0
3  604000.0          4         3.00         1960     5000     1.0           0
4  510000.0          3         2.00         1680     8080     1.0           0

      condition  grade  sqft_above  ...  east  fareast  north  west  \
0             3      7         1180  ...    0         0      0      1
1             3      7         2170  ...    0         0      0      1
2             3      6          770  ...    1         0      0      0
3             5      7         1050  ...    0         0      0      1
4             3      8         1680  ...    0         1      0      0

      has_basement  log_sqft_living  log_sqft_living15  log_sqft_above  \
0                0         7.073270         7.200425         7.073270
1                1         7.851661         7.432484         7.682482
2                0         6.646391         7.908387         6.646391
3                1         7.580700         7.215240         6.956545
4                0         7.426549         7.495542         7.426549

      log_sqft_lot  log_sqft_lot15
0         8.639411         8.639411
1         8.887653         8.941022
2         9.210340         8.994917
3         8.517193         8.517193
4         8.997147         8.923058

[5 rows x 38 columns]
```

```
[105]: data.columns
```

```
[105]: Index(['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors',
      'waterfront', 'condition', 'grade', 'sqft_above', 'lat', 'long',
      'sqft_living15', 'sqft_lot15', 'log_price', 'has_view', 'april',
      'august', 'december', 'february', 'july', 'june', 'march', 'may',
      'november', 'october', 'september', 'age<30', 'east', 'fareast',
      'north', 'west', 'has_basement', 'log_sqft_living', 'log_sqft_living15',
      'log_sqft_above', 'log_sqft_lot', 'log_sqft_lot15'],
      dtype='object')
```

### 0.9.1 HEATMAP


```
[106]: variables = data[['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', '
↳ 'floors',
        'waterfront', 'condition', 'grade', 'sqft_above',
        'lat', 'long', 'sqft_living15', 'sqft_lot15', 'has_view',
        'age<30', 'east', 'fareast',
        'north', 'west', 'has_basement']]

corr = variables.corr().abs()

fig, ax=plt.subplots(figsize=(14,14))
matrix = np.triu(corr) # Getting the Upper Triangle of the correlation matrix
cbar_kws={"label": "Correlation", "shrink":0.8}
heatmap = sns.heatmap(data = corr, cmap='Blues', linewidths = 1, square= True, 
↳ ax=ax, annot=True, mask=matrix, fmt= ".2f", cbar_kws=cbar_kws)
fig.suptitle('Heatmap of Correlation Between All Variables (Including Target)', 
↳ fontsize=20, y=.84, x = .43, fontname='Arial');
heatmap;

fig.savefig('./images/Heatmap_Correlation.png', dpi=300);
```

Heatmap of Correlation Between All Variables (Including Target)



Write a function to show the most correlated pairs:

```
[107]: def show_corr_pairs(data):
    dataCorr = data.corr().abs()
    dataCorr = dataCorr.mask(np.triu(np.ones(dataCorr.shape)).astype(np.bool))
    # convert upper triangle of values to NaN to remove repeated values from the
    # table
    dataCorr = dataCorr.stack().reset_index().sort_values(0, ascending=False)
    # 0 is the column automatically generated by the stacking
    dataCorr = dataCorr[(dataCorr[0]>.7) & (dataCorr[0]<1)]
    dataCorr = dataCorr.rename(columns = {'level_0': 'var1', 'level_1': 'var2', 0: 'corrcoef'})
    return dataCorr
```

```
[108]: show_corr_pairs(variables)
```

```
[108]:
```

	var1	var2	corrcoef
39	sqft_above	sqft_living	0.867865
82	sqft_lot15	sqft_lot	0.832800
69	sqft_living15	sqft_living	0.763589
31	grade	sqft_living	0.755051
44	sqft_above	grade	0.747496
5	sqft_living	bathrooms	0.743058
75	sqft_living15	sqft_above	0.735825
74	sqft_living15	grade	0.711495

- sqft\_living correlates highly with sqft\_above and sqft\_living.
- sqft\_living correlates highly with grade and bathrooms too.
- sqft\_lot15 correlates highly with sqft\_lot.

```
[109]: variables.corr()['price'].map(abs).sort_values(ascending=False)
```

```
[109]: price          1.000000
sqft_living    0.685533
grade          0.675277
sqft_living15  0.600749
sqft_above     0.586275
bathrooms      0.508569
has_view       0.355573
lat            0.344114
bedrooms       0.321149
east           0.270383
floors         0.264440
waterfront     0.191616
has_basement   0.182308
age<30         0.152268
sqft_lot       0.133435
sqft_lot15     0.122568
fareast        0.100217
west           0.074890
north          0.045800
condition      0.043828
long           0.015586
Name: price, dtype: float64
```

- sqft\_living seems to have the greatest correlation with price.

```
[110]: df_corr = abs(variables.corr()) > 0.7
df_corr.sum()
```

```
[110]: price          1
bedrooms          1
bathrooms         2
sqft_living       5
```

```

sqft_lot      2
floors        1
waterfront    1
condition     1
grade         4
sqft_above    4
lat           1
long          1
sqft_living15 4
sqft_lot15    2
has_view      1
age<30        1
east          1
fareast       1
north         1
west          1
has_basement  1
dtype: int64

```

- Square Foot Living is the variable with the greatest collinearity to others too.

### 0.9.2 Take away from the Heat Map

- `sqft_living`, `sqft_above`, `sqft_living15` correlate highly. Keep `sqft_living` as it correlates with price the highest.
- `grade` and `bathrooms` also correlate highly with `sqft_living`. But let's keep these variables since they give a different type of information.
- Do not use `lat` and `long` since they are redundant with location variables.
- `sqft_lot15` and `sqft_lot` correlate highly. Keep `sqft_lot` as it correlates with price a bit more.

## 0.10 Regression Assumptions Check Functions:

### 0.10.1 Linearity:

- There should be a linear relationship between the response/target variable and predictors.
- Check for this once using scatterplots and then visually inspecting the scatterplots for linearity..

```
[111]: data.columns
```


```

[111]: Index(['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors',
            'waterfront', 'condition', 'grade', 'sqft_above', 'lat', 'long',
            'sqft_living15', 'sqft_lot15', 'log_price', 'has_view', 'april',
            'august', 'december', 'february', 'july', 'june', 'march', 'may',
            'november', 'october', 'september', 'age<30', 'east', 'fareast',
            'north', 'west', 'has_basement', 'log_sqft_living', 'log_sqft_living15',
            'log_sqft_above', 'log_sqft_lot', 'log_sqft_lot15'],
            dtype='object')


```


```
[112]: # Linearity against `price`:
continuous = ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors',
              'condition', 'grade']


for column in continuous:
    sns.jointplot(x=column, y="price", data=data, kind='reg',
                  label=column, joint_kws={'line_kws':{'color':'red'}})
    plt.legend()
    plt.show()
```






















- Looks like a linear relationship except `sqft_lot`

```
[113]: # Linearity against `log_price`:


continuous = ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors',
              'condition', 'grade']


for column in continuous:
    sns.jointplot(x=column, y="log_price", data=data, kind='reg',
                  label=column, joint_kws={'line_kws':{'color':'red'}})
    plt.legend()
    plt.show()
```




















- Looks like a linear relationship for all variables.

### 0.10.2 Check for Normality and Homoscedasticity:

- The model residuals should follow a normal distribution.
- The residuals should be evenly spread through range (The variance of residual is uniform).

```
[114]: def normality_homoscedasticity(model):
    fig, ((ax1, ax2, ax3)) = plt.subplots(ncols=3, figsize=(16, 4))

    ax1.hist(model.resid_pearson, bins=20, edgecolor='k')
    ax1.set_xlabel("Normalized residuals", fontsize=14)
    ax1.set_ylabel("Count", fontsize=14)
    ax1.set_title("Histogram of normalized residuals (NORMALITY)", fontsize =11)
```

```

from statsmodels.graphics.gofplots import qqplot
qqplot(model.resid_pearson, line='45', fit='True', ax = ax2, alpha=0.8,
↪markerfacecolor='#1f77b4')
ax2.set_xlabel("Theoretical quantiles",fontsize=14)
ax2.set_ylabel("Ordered Values",fontsize=14)
ax2.set_title("Q-Q plot of normalized residuals (NORMALITY)", fontsize =11)

ax3.scatter(x=model.fittedvalues, y=model.resid)
xmin=min(model.fittedvalues)
xmax = max(model.fittedvalues)
plt.hlines(y=0,xmin=xmin,xmax=xmax,color='red',linestyle='--',lw=3)
ax3.set_xlabel("Fitted values",fontsize=14)
ax3.set_ylabel("Residuals",fontsize=14)
ax3.set_title("Fitted vs. residuals plot (HOMOSCADECITY)", fontsize =11)
plt.grid(True)

return(ax)

```

### 0.10.3 Check for absence of multicollinearity:

- Multicollinearity occurs when 2 or more of the independent variables are highly correlated with each other.
- VIF (variance inflation factor) is a measure for the increase of the variance of the parameter estimates if an additional variable is added to the linear regression.
- If VIF is greater than 5, then the explanatory variable is highly collinear with another explanatory variable.

```

[115]: #https://github.com/cwf231/linear_regression_guided_practice
from statsmodels.stats.outliers_influence import variance_inflation_factor

def create_vif_dictionary(X):
    vif_dct = {}
    for i in range(len(X.columns)): # Loop through each row and set the
↪variable name to the VIF.
        vif = variance_inflation_factor(X.values, i) # Calculate VIF
        v = X.columns[i] # Extract column name for dictionary key.
        vif_dct[v] = vif # Set value in dictionary.
    return vif_dct

def multicollinearity(X):
    multicollinearity = pd.DataFrame(create_vif_dictionary(X), index=[0]).T
    return multicollinearity.sort_values(by = 0, ascending =False).apply(lambda
↪x: x.apply('{0:.4f}'.format))

```

## 0.11 Regression Modeling:

### 0.11.1 BASELINE MODEL #1

- The baseline model is using the most highly correlated variable with price: `sqft_living`

```
[116]: y = data['price']
X = data['sqft_living']

X.shape, y.shape
```

```
[116]: ((20904,), (20904,))
```

```
[117]: X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
model.summary()
```

```
[117]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                        OLS Regression Results
=====
Dep. Variable:          price      R-squared:                0.470
Model:                  OLS       Adj. R-squared:            0.470
Method:                 Least Squares   F-statistic:            1.853e+04
Date:                  Tue, 23 Aug 2022   Prob (F-statistic):      0.00
Time:                  20:18:54    Log-Likelihood:         -2.8804e+05
No. Observations:      20904        AIC:                   5.761e+05
Df Residuals:          20902        BIC:                   5.761e+05
Df Model:              1
Covariance Type:       nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
const          1.031e+04    4143.504      2.489     0.013     2191.002     1.84e+04
sqft_living    252.8318      1.857     136.134     0.000      249.192      256.472
=====
Omnibus:            8006.855   Durbin-Watson:           1.989
Prob(Omnibus):      0.000   Jarque-Bera (JB):        52765.454
Skew:               1.698   Prob(JB):                0.00
Kurtosis:           10.004   Cond. No.                5.73e+03
=====
```


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.73e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
"""
```

```
[118]: normality_homoscedasticity(model);
```



### Violation of normality and homoscedasticity:

- The Distribution of the residuals are NOT normal. The blue dots are the observed data while the red regression line is the prediction on the second graph. The residuals are NOT normally distributed as the blue dots are not falling on the red line. We may fix this by transforming the target variable and/or independent variables.
- There seems like a violation of Homoskedasticity as well since the dots around the red line are not symmetric and follow a cone-like shape. We can try log transforming the target variable.

#### 0.11.2 MODEL #2

- Using log transformed `log_price` as the target variable.

```
[119]: y = data['log_price']
X = data['sqft_living']
X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
model.summary()
```

```
[119]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  log_price    R-squared:                0.463
Model:                            OLS       Adj. R-squared:            0.463
Method:                 Least Squares    F-statistic:                1.802e+04
Date:                Tue, 23 Aug 2022    Prob (F-statistic):          0.00
Time:                        20:18:55    Log-Likelihood:            -9171.7
No. Observations:                20904    AIC:                       1.835e+04
Df Residuals:                    20902    BIC:                       1.836e+04
Df Model:                            1
Covariance Type:                nonrobust
=====
```



	coef	std err	t	P> t	[0.025	0.975]
-----	-----	-----	-----	-----	-----	-----
const	12.2174	0.007	1833.395	0.000	12.204	12.230
sqft_living	0.0004	2.99e-06	134.225	0.000	0.000	0.000
=====	=====	=====	=====	=====	=====	=====
Omnibus:	76.821		Durbin-Watson:		1.989	
Prob(Omnibus):	0.000		Jarque-Bera (JB):		61.899	
Skew:	0.057		Prob(JB):		3.62e-14	
Kurtosis:	2.759		Cond. No.		5.73e+03	
=====	=====	=====	=====	=====	=====	=====


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.73e+03. This might indicate that there are strong multicollinearity or other numerical problems.

"""

```
[120]: normality_homoscedecity(model);
```



- Normality and homoscedecity seem to be restored when we used log-transformed target variable!

### Summary interpretation:

- p value for `sqft_living` is statistically significant, meaning we can reject the null hypothesis that `sqft_living` does not correlate with price. In other words we can also say: there is enough evidence in favor of the idea that change in `sqft_living` is associated with change in `price` at the population level. This variable is a worthwhile addition to our regression model.
- A Coefficient of Determination **R-Squared** value of .46 means that 46% of the variability in price is explained by `sqft_living`. R-Squared explains how good our model is when compared to a baseline model where  $y = mx$ . We need to boost this number.

### 0.11.3 MODEL #3


- Using log transformed `log_sqft_living` as the predictor variable to see if it would improve `R2`.

```
[121]: y = data['log_price']
X = data['log_sqft_living']
X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
model.summary()
```

```
[121]: <class 'statsmodels.iolib.summary.Summary'>
"""
                                OLS Regression Results
=====
Dep. Variable:                  log_price    R-squared:                  0.437
Model:                            OLS      Adj. R-squared:            0.437
Method:                 Least Squares    F-statistic:                1.625e+04
Date:                Tue, 23 Aug 2022    Prob (F-statistic):          0.00
Time:                  20:18:55          Log-Likelihood:            -9658.4
No. Observations:                20904    AIC:                      1.932e+04
Df Residuals:                  20902    BIC:                      1.934e+04
Df Model:                            1
Covariance Type:                nonrobust
=====
===
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
---
const                6.9135      0.048    143.586      0.000      6.819
7.008
log_sqft_living      0.8124      0.006    127.459      0.000      0.800
0.825
=====
Omnibus:                 162.613    Durbin-Watson:             1.991
Prob(Omnibus):            0.000    Jarque-Bera (JB):          126.560
Skew:                     0.107    Prob(JB):                  3.29e-28
Kurtosis:                 2.685    Cond. No.                  139.
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
"""
```

```
[122]: normality_homoscedasticity(model);
```



- This was worse in terms of R2 and even normality. Let's go back to using un-transformed `sqft_living`.

#### 0.11.4 MODEL #4

- Using `sqft_living` and `sqft_lot` as the 2 basic area variables.

```
[123]: y = data['log_price']
X = data[['sqft_living', 'sqft_lot']]
X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
model.summary()
```

```
[123]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  log_price    R-squared:                  0.466
Model:                            OLS      Adj. R-squared:              0.465
Method:                 Least Squares    F-statistic:                  9102.
Date:                Tue, 23 Aug 2022    Prob (F-statistic):           0.00
Time:                  20:18:56          Log-Likelihood:              -9121.3
No. Observations:                20904    AIC:                        1.825e+04
Df Residuals:                    20901    BIC:                        1.827e+04
Df Model:                          2
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	12.2236	0.007	1830.876	0.000	12.210	12.237
sqft_living	0.0004	3.09e-06	132.479	0.000	0.000	0.000
sqft_lot	-2.175e-06	2.16e-07	-10.049	0.000	-2.6e-06	-1.75e-06

```

=====
Omnibus:                        65.587    Durbin-Watson:                1.989
Prob(Omnibus):                  0.000    Jarque-Bera (JB):              53.110
Skew:                           0.047    Prob(JB):                      2.93e-12

```

Kurtosis: 2.771 Cond. No. 4.21e+04  
=====


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.21e+04. This might indicate that there are strong multicollinearity or other numerical problems.

"""

```
[124]: normality_homoscedecity(model);
```



```
[125]: multicollinearity(X)
```

```
[125]:          0
const      6.6484
sqft_living 1.0737
sqft_lot    1.0737
```

```
[126]: print(data.corr()['price']['sqft_living'])
print(data.corr()['price']['sqft_lot'])
```

```
0.6855331909828829
0.13343455593875633
```

```
[127]: with plt.style.context('seaborn-talk'):
fig, (ax1, ax2) = plt.subplots(ncols=2, nrows=1, figsize=(12, 5))
fig.set_tight_layout(True)

sns.regplot(x="sqft_living", y="price", ax=ax1, data=data,
↳line_kws={"color": "orange", "label": "R^2=.69"})
ax1.legend()
ax1.yaxis.set_major_formatter(formatter)
ax1.set_title('Relationship Between Total Square Footage and_
↳Price', fontsize=14)
```


```

ax1.set_xlabel("Square Footage of the House",fontsize=12)
ax1.set_ylabel("Price",fontsize=12)

sns.regplot(x="sqft_lot", y="price", ax=ax2, data=data, line_kws={"color": "orange", "label": "R^2=.13"})
ax2.legend()
ax2.yaxis.set_major_formatter(formatter)
ax2.set_title('Relationship Between Lot Square Footage and Price',fontsize=14)
ax2.set_xlabel("Square Footage of the Lot",fontsize=12)
ax2.set_ylabel("Price",fontsize=12)

fig.savefig('./images/sqft_living_sqft_lot.png', dpi=300);

```



- `sqft_lot` adds little to the model increasing R-squared slightly from 0.463 to 0.466.
- But this variable is still statistically significant, so it is still worthwhile adding it to the model.

### 0.11.5 MODEL #5

- Adding other meaningful variables except age, season and location

```

[128]: variables = ['sqft_living', 'sqft_lot', 'condition', 'grade']

y = data['log_price']
X = data[variables]

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
model.rsquared

```

[128]: 0.552871437510315

```
[129]: variables = ['sqft_living', 'sqft_lot', 'condition', 'grade', 'has_view',
↳ 'waterfront']

y = data['log_price']
X = data[variables]

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
model.rsquared
```

[129]: 0.5733582227842691

```
[130]: variables = ['sqft_living', 'sqft_lot', 'condition', 'grade', 'has_view',
↳ 'waterfront',
↳ 'bedrooms', 'bathrooms']

y = data['log_price']
X = data[variables]

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
model.rsquared
```

[130]: 0.5742774950797981

```
[131]: variables = ['sqft_living', 'sqft_lot', 'condition', 'grade', 'has_view',
↳ 'waterfront',
↳ 'bedrooms', 'bathrooms', 'has_basement']

y = data['log_price']
X = data[variables]

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
print(model.rsquared)
```

0.5813996239007405

```
[132]: variables = ['sqft_living', 'sqft_lot', 'condition', 'grade', 'has_view',
↳ 'waterfront',
↳ 'bedrooms', 'bathrooms', 'has_basement', 'floors']

y = data['log_price']
X = data[variables]

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
```

```
print(model.rsquared)
```

```
0.5839041800981312
```

```
[133]: model.summary()
```

```
[133]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

#### OLS Regression Results

```
=====
Dep. Variable:          log_price    R-squared:                0.584
Model:                  OLS          Adj. R-squared:            0.584
Method:                 Least Squares    F-statistic:            2932.
Date:                  Tue, 23 Aug 2022    Prob (F-statistic):      0.00
Time:                  20:19:00          Log-Likelihood:         -6504.1
No. Observations:      20904            AIC:                   1.303e+04
Df Residuals:          20893            BIC:                   1.312e+04
Df Model:              10
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	10.8146	0.026	417.754	0.000	10.764	10.865
sqft_living	0.0002	5.43e-06	38.001	0.000	0.000	0.000
sqft_lot	-1.471e-06	1.98e-07	-7.413	0.000	-1.86e-06	-1.08e-06
condition	0.0954	0.004	25.884	0.000	0.088	0.103
grade	0.1911	0.003	57.642	0.000	0.185	0.198
has_view	0.1825	0.008	21.539	0.000	0.166	0.199
waterfront	0.4293	0.031	13.709	0.000	0.368	0.491
bedrooms	-0.0201	0.003	-5.965	0.000	-0.027	-0.013
bathrooms	-0.0279	0.005	-5.386	0.000	-0.038	-0.018
has_basement	0.1159	0.005	21.687	0.000	0.105	0.126
floors	0.0636	0.006	11.214	0.000	0.052	0.075

```
=====
Omnibus:                11.677    Durbin-Watson:           1.980
Prob(Omnibus):          0.003    Jarque-Bera (JB):        11.672
Skew:                   0.058    Prob(JB):                0.00292
Kurtosis:               3.011    Cond. No.                 2.24e+05
=====
```


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.24e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
"""
```

```
[134]: normality_homoscedacity(model);
```



```
[135]: multicollinearity(X)
```

```
[135]:
```

	0
const	128.3504
sqft_living	4.2572
bathrooms	2.8699
grade	2.7510
floors	1.7974
bedrooms	1.7152
has_basement	1.2957
has_view	1.1752
sqft_lot	1.1586
condition	1.0991
waterfront	1.0628

- We increased R2 to 0.584 while still keeping normality and homoscedacity intact.
- All variables are statistically significant.

### 0.11.6 MODEL #6

- Adding the month the house was sold (January is reference) to the model.

```
[136]: variables = ['sqft_living', 'sqft_lot', 'condition', 'grade', 'has_view', 'waterfront',
                    'bedrooms', 'bathrooms', 'has_basement', 'floors',
                    'april', 'august', 'december', 'february', 'july', 'june', 'march',
                    'may', 'november', 'october', 'september']

y = data['log_price']
X = data[variables]

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
```



```
model.summary()
```

```
[136]: <class 'statsmodels.iolib.summary.Summary'>
```

```
"""
                                OLS Regression Results
=====
Dep. Variable:                log_price    R-squared:                0.586
Model:                        OLS         Adj. R-squared:           0.586
Method:                       Least Squares   F-statistic:             1409.
Date:                         Tue, 23 Aug 2022   Prob (F-statistic):       0.00
Time:                         20:19:00        Log-Likelihood:          -6446.0
No. Observations:             20904          AIC:                    1.294e+04
Df Residuals:                 20882          BIC:                    1.311e+04
Df Model:                     21
Covariance Type:              nonrobust
=====
                                coef    std err          t      P>|t|      [0.025    0.975]
-----
const                10.7851      0.028    389.308      0.000     10.731    10.839
sqft_living           0.0002    5.41e-06    38.221      0.000      0.000      0.000
sqft_lot            -1.469e-06    1.98e-07   -7.424      0.000   -1.86e-06   -1.08e-06
condition            0.0964      0.004    26.187      0.000      0.089      0.104
grade               0.1908      0.003    57.690      0.000      0.184      0.197
has_view            0.1817      0.008    21.502      0.000      0.165      0.198
waterfront          0.4284      0.031    13.714      0.000      0.367      0.490
bedrooms            -0.0204      0.003     -6.085      0.000     -0.027     -0.014
bathrooms           -0.0277      0.005     -5.363      0.000     -0.038     -0.018
has_basement        0.1155      0.005    21.662      0.000      0.105      0.126
floors              0.0639      0.006    11.294      0.000      0.053      0.075
april               0.0813      0.013      6.341      0.000      0.056      0.106
august              0.0143      0.013      1.088      0.277     -0.011      0.040
december            -0.0041      0.014     -0.293      0.769     -0.031      0.023
february           0.0049      0.014      0.342      0.733     -0.023      0.033
july               0.0178      0.013      1.380      0.168     -0.007      0.043
june               0.0272      0.013      2.104      0.035      0.002      0.052
march              0.0571      0.013      4.324      0.000      0.031      0.083
may                0.0390      0.013      3.067      0.002      0.014      0.064
november           0.0033      0.014      0.236      0.813     -0.024      0.031
october            0.0191      0.013      1.442      0.149     -0.007      0.045
september          0.0160      0.013      1.200      0.230     -0.010      0.042
=====
Omnibus:                 9.523    Durbin-Watson:           1.983
Prob(Omnibus):           0.009    Jarque-Bera (JB):        9.509
Skew:                    0.052    Prob(JB):                0.00861
Kurtosis:                3.015    Cond. No.                2.78e+05
=====
```


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large,  $2.78e+05$ . This might indicate that there are strong multicollinearity or other numerical problems.

"""

```
[137]: normality_homoscedasticity(model);
```



```
[138]: multicollinearity(X)
```

```
[138]:
```

	0
const	147.7278
sqft_living	4.2592
may	3.0693
april	2.9614
july	2.9319
june	2.9041
bathrooms	2.8723
grade	2.7531
august	2.7078
march	2.6760
october	2.6552
september	2.5760
december	2.3326
november	2.2822
february	2.1584
floors	1.7988
bedrooms	1.7161
has_basement	1.2968
has_view	1.1760
sqft_lot	1.1589
condition	1.1038
waterfront	1.0632

- R2 increased very slightly from 0.584 to .586 - normality and homoscedasticity are intact, and multicollinearity is absent.
- It seems like it is the months of spring that have an impact on price. But there are also a lot of non-significant months. Let's remove the nonsignificant months from the model to increase its performance (non-significance means that there is insufficient evidence in your sample to conclude that a correlation exists).

#### Removing nonsignificant months:

```
[139]: variables = ['sqft_living', 'log_sqft_lot', 'condition', 'grade', 'has_view',
↳ 'waterfront',
        'bedrooms', 'bathrooms', 'has_basement', 'floors',
        'april', 'march', 'may', 'june']

y = data['log_price']
X = data[variables]

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
model.summary()
```

```
[139]: <class 'statsmodels.iolib.summary.Summary'>
```

```
"""
                                OLS Regression Results
=====
Dep. Variable:                  log_price    R-squared:                  0.591
Model:                            OLS       Adj. R-squared:            0.591
Method:                 Least Squares    F-statistic:                 2157.
Date:                 Tue, 23 Aug 2022    Prob (F-statistic):          0.00
Time:                 20:19:01           Log-Likelihood:             -6322.2
No. Observations:                20904    AIC:                       1.267e+04
Df Residuals:                    20889    BIC:                       1.279e+04
Df Model:                        14
Covariance Type:                nonrobust
=====
                                coef    std err          t      P>|t|      [0.025    0.975]
-----
const                11.3883      0.043    264.906      0.000      11.304     11.473
sqft_living           0.0002   5.66e-06     41.670      0.000       0.000       0.000
log_sqft_lot        -0.0665      0.004    -17.676      0.000      -0.074     -0.059
condition            0.0979      0.004     26.765      0.000       0.091       0.105
grade                0.1904      0.003     57.923      0.000       0.184       0.197
has_view             0.1827      0.008     21.762      0.000       0.166       0.199
waterfront           0.4553      0.031     14.648      0.000       0.394       0.516
bedrooms            -0.0171      0.003     -5.148      0.000      -0.024     -0.011
bathrooms           -0.0331      0.005     -6.432      0.000      -0.043     -0.023
has_basement         0.0890      0.006     16.043      0.000       0.078       0.100
floors               0.0194      0.006      3.097      0.002       0.007       0.032
=====
```

april	0.0703	0.008	9.248	0.000	0.055	0.085
march	0.0462	0.008	5.626	0.000	0.030	0.062
may	0.0269	0.007	3.633	0.000	0.012	0.041
june	0.0181	0.008	2.349	0.019	0.003	0.033

```
=====
Omnibus:                9.095    Durbin-Watson:                1.979
Prob(Omnibus):          0.011    Jarque-Bera (JB):          9.121
Skew:                   0.045    Prob(JB):                  0.0105
Kurtosis:               3.049    Cond. No.                  4.27e+04
=====
```


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.27e+04. This might indicate that there are strong multicollinearity or other numerical problems.

"""

```
[140]: normality_homoscedecity(model);
```



```
[141]: multicollinearity(X)
```

```
[141]:
```

	0
const	360.1094
sqft_living	4.7209
bathrooms	2.8832
grade	2.7509
floors	2.2386
bedrooms	1.7102
log_sqft_lot	1.6723
has_basement	1.4218
has_view	1.1746
condition	1.1024
waterfront	1.0657
may	1.0548

april	1.0535
june	1.0526
march	1.0483

- R2 increased from 0.584 to 0.591 with the addition of month variable.
- Normality and homoscedasticity intact.
- All variables are statistically significant.

### 0.11.7 MODEL #7

- adding age<30

```
[142]: variables = ['sqft_living', 'log_sqft_lot', 'condition', 'grade', 'has_view', 'waterfront',
                  'bedrooms', 'bathrooms', 'has_basement', 'floors',
                  'april', 'march', 'may', 'june',
                  'age<30']

y = data['log_price']
X = data[variables]

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
model.summary()
```

```
[142]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  log_price    R-squared:                0.602
Model:                            OLS       Adj. R-squared:            0.601
Method:                 Least Squares    F-statistic:                2104.
Date:                Tue, 23 Aug 2022    Prob (F-statistic):          0.00
Time:                  20:19:02    Log-Likelihood:             -6045.4
No. Observations:                20904    AIC:                        1.212e+04
Df Residuals:                    20888    BIC:                        1.225e+04
Df Model:                          15
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	11.4741	0.043	269.472	0.000	11.391	11.558
sqft_living	0.0002	5.59e-06	42.195	0.000	0.000	0.000
log_sqft_lot	-0.0762	0.004	-20.405	0.000	-0.083	-0.069
condition	0.0684	0.004	17.914	0.000	0.061	0.076
grade	0.1988	0.003	60.928	0.000	0.192	0.205
has_view	0.1629	0.008	19.556	0.000	0.147	0.179
waterfront	0.4487	0.031	14.626	0.000	0.389	0.509

bedrooms	-0.0285	0.003	-8.570	0.000	-0.035	-0.022
bathrooms	0.0065	0.005	1.223	0.221	-0.004	0.017
has_basement	0.0657	0.006	11.809	0.000	0.055	0.077
floors	0.0656	0.006	10.114	0.000	0.053	0.078
april	0.0696	0.008	9.272	0.000	0.055	0.084
march	0.0456	0.008	5.627	0.000	0.030	0.061
may	0.0302	0.007	4.132	0.000	0.016	0.045
june	0.0178	0.008	2.339	0.019	0.003	0.033
age<30	-0.1664	0.007	-23.677	0.000	-0.180	-0.153

```
=====
Omnibus:                17.606    Durbin-Watson:                1.987
Prob(Omnibus):          0.000    Jarque-Bera (JB):            20.068
Skew:                   -0.003    Prob(JB):                    4.39e-05
Kurtosis:               3.152    Cond. No.:                   4.28e+04
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.28e+04. This might indicate that there are strong multicollinearity or other numerical problems.

"""

**Remove bathrooms as it is not significant any more:**

- It also correlates highly with `sqft_living`.

```
[143]: variables = ['sqft_living', 'log_sqft_lot', 'condition', 'grade', 'has_view',
↳ 'waterfront',
           'bedrooms', 'has_basement', 'floors',
           'april', 'march', 'may', 'june',
           'age<30']

y = data['log_price']
X = data[variables]

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
model.summary()
```

```
[143]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                log_price    R-squared:                0.602
Model:                        OLS         Adj. R-squared:            0.601
Method:                       Least Squares    F-statistic:            2254.
```

```

Date:                Tue, 23 Aug 2022    Prob (F-statistic):        0.00
Time:                20:19:02           Log-Likelihood:          -6046.1
No. Observations:    20904             AIC:                    1.212e+04
Df Residuals:        20889             BIC:                    1.224e+04
Df Model:            14
Covariance Type:     nonrobust

```

	coef	std err	t	P> t	[0.025	0.975]
const	11.4737	0.043	269.467	0.000	11.390	11.557
sqft_living	0.0002	5.26e-06	45.307	0.000	0.000	0.000
log_sqft_lot	-0.0763	0.004	-20.467	0.000	-0.084	-0.069
condition	0.0687	0.004	18.039	0.000	0.061	0.076
grade	0.1992	0.003	61.375	0.000	0.193	0.206
has_view	0.1630	0.008	19.567	0.000	0.147	0.179
waterfront	0.4489	0.031	14.632	0.000	0.389	0.509
bedrooms	-0.0275	0.003	-8.516	0.000	-0.034	-0.021
has_basement	0.0670	0.005	12.278	0.000	0.056	0.078
floors	0.0670	0.006	10.472	0.000	0.054	0.080
april	0.0696	0.008	9.270	0.000	0.055	0.084
march	0.0455	0.008	5.619	0.000	0.030	0.061
may	0.0301	0.007	4.122	0.000	0.016	0.044
june	0.0179	0.008	2.353	0.019	0.003	0.033
age<30	-0.1638	0.007	-24.527	0.000	-0.177	-0.151
Omnibus:	17.640		Durbin-Watson:	1.987		
Prob(Omnibus):	0.000		Jarque-Bera (JB):	20.104		
Skew:	-0.004		Prob(JB):	4.31e-05		
Kurtosis:	3.152		Cond. No.	4.28e+04		


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.28e+04. This might indicate that there are strong multicollinearity or other numerical problems.

""

```
[144]: normality_homoscedasticity(model);
```



```
[145]: multicollinearity(X)
```

```
[145]:
```

	0
const	362.7151
sqft_living	4.1756
grade	2.7550
floors	2.3907
age<30	2.1350
log_sqft_lot	1.6902
bedrooms	1.6554
has_basement	1.4128
condition	1.2280
has_view	1.1866
waterfront	1.0657
may	1.0551
april	1.0535
june	1.0524
march	1.0483

- R2 increased from .591 to .602 with the addition of **age** variable.
- Normality and homoscedasticity intact.

### 0.11.8 FINAL MODEL #8

- adding location variables

```
[146]: variables = ['sqft_living', 'sqft_lot', 'condition', 'grade', 'has_view',
    ↪, 'waterfront',
    'bedrooms', 'has_basement', 'floors',
    'april', 'march', 'may', 'june',
    'age<30',
    'east', 'fareast', 'north', 'west']

y = data['log_price']
X = data[variables]
```



```
X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
model.summary()
```

[146]: <class 'statsmodels.iolib.summary.Summary'>

```
"""
                                OLS Regression Results
=====
Dep. Variable:                  log_price    R-squared:                  0.753
Model:                            OLS      Adj. R-squared:              0.752
Method:                 Least Squares    F-statistic:                 3531.
Date:                 Tue, 23 Aug 2022    Prob (F-statistic):          0.00
Time:                      20:19:02      Log-Likelihood:             -1066.0
No. Observations:          20904         AIC:                        2170.
Df Residuals:              20885         BIC:                        2321.
Df Model:                   18
Covariance Type:            nonrobust
=====
                                coef    std err          t      P>|t|      [0.025    0.975]
-----
const                10.6461      0.020    527.661      0.000      10.607     10.686
sqft_living           0.0002    3.98e-06    54.679      0.000      0.000      0.000
sqft_lot             1.462e-06    1.58e-07     9.252      0.000    1.15e-06    1.77e-06
condition            0.0834      0.003    27.687      0.000      0.077      0.089
grade               0.1608      0.003    62.077      0.000      0.156      0.166
has_view            0.1427      0.007    21.628      0.000      0.130      0.156
waterfront          0.5171      0.024    21.397      0.000      0.470      0.564
bedrooms            -0.0030      0.003    -1.155      0.248     -0.008      0.002
has_basement        0.0173      0.004     4.032      0.000      0.009      0.026
floors              0.0303      0.005     6.315      0.000      0.021      0.040
april               0.0677      0.006    11.452      0.000      0.056      0.079
march              0.0561      0.006     8.786      0.000      0.044      0.069
may                0.0169      0.006     2.939      0.003      0.006      0.028
june               0.0090      0.006     1.494      0.135     -0.003      0.021
age<30             -0.0246      0.005    -4.533      0.000     -0.035     -0.014
east               0.4867      0.005    90.505      0.000      0.476      0.497
fareast            0.4006      0.006    62.555      0.000      0.388      0.413
north              0.3322      0.008    40.583      0.000      0.316      0.348
west              0.5308      0.005   104.297      0.000      0.521      0.541
=====
Omnibus:                 581.185    Durbin-Watson:              1.989
Prob(Omnibus):           0.000    Jarque-Bera (JB):          1409.328
Skew:                   -0.090    Prob(JB):                  9.30e-307
Kurtosis:                4.259    Cond. No.                  2.25e+05
=====
```

Notes:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
[2] The condition number is large, 2.25e+05. This might indicate that there are
strong multicollinearity or other numerical problems.
"""
```

Remove bedrooms and june from the model:

```
[147]: variables = variables = ['sqft_living', 'sqft_lot', 'condition', 'grade',
    ↪ 'has_view', 'waterfront',
    'has_basement', 'floors',
    'april', 'march', 'may',
    'age<30',
    'east', 'fareast', 'north', 'west']

y = data['log_price']
X = data[variables]

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
model.summary()
```

```
[147]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                        OLS Regression Results
=====
Dep. Variable:          log_price    R-squared:                0.753
Model:                  OLS         Adj. R-squared:            0.752
Method:                 Least Squares   F-statistic:              3972.
Date:                  Tue, 23 Aug 2022   Prob (F-statistic):       0.00
Time:                  20:19:02         Log-Likelihood:           -1067.8
No. Observations:      20904           AIC:                     2170.
Df Residuals:          20887           BIC:                     2305.
Df Model:              16
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	10.6384	0.019	563.259	0.000	10.601	10.675
sqft_living	0.0002	3.43e-06	62.667	0.000	0.000	0.000
sqft_lot	1.48e-06	1.57e-07	9.407	0.000	1.17e-06	1.79e-06
condition	0.0834	0.003	27.712	0.000	0.078	0.089
grade	0.1612	0.003	62.788	0.000	0.156	0.166
has_view	0.1433	0.007	21.805	0.000	0.130	0.156
waterfront	0.5184	0.024	21.474	0.000	0.471	0.566
has_basement	0.0171	0.004	4.005	0.000	0.009	0.026
floors	0.0300	0.005	6.257	0.000	0.021	0.039

april	0.0664	0.006	11.346	0.000	0.055	0.078
march	0.0547	0.006	8.651	0.000	0.042	0.067
may	0.0156	0.006	2.742	0.006	0.004	0.027
age<30	-0.0240	0.005	-4.440	0.000	-0.035	-0.013
east	0.4868	0.005	90.527	0.000	0.476	0.497
fareast	0.4012	0.006	62.808	0.000	0.389	0.414
north	0.3324	0.008	40.612	0.000	0.316	0.348
west	0.5316	0.005	105.152	0.000	0.522	0.541

```
=====
Omnibus:                    579.603    Durbin-Watson:                1.988
Prob(Omnibus):              0.000    Jarque-Bera (JB):             1406.417
Skew:                      -0.088    Prob(JB):                     3.99e-306
Kurtosis:                   4.258    Cond. No.                     2.24e+05
=====
```


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.24e+05. This might indicate that there are strong multicollinearity or other numerical problems.

"""

```
[148]: normality_homoscedacity(model);
```



```
[149]: multicollinearity(X)
```

```
[149]:
const          0
sqft_living    114.8988
grade          2.8630
age<30         2.7764
age<30         2.2580
floors         2.1560
west           1.8748
east           1.5564
has_basement   1.3967
```

fareast	1.3886
condition	1.2346
sqft_lot	1.2261
has_view	1.1903
north	1.1738
waterfront	1.0626
may	1.0318
april	1.0313
march	1.0295

- R2 increased from .602 to .753 with the addition of locations!
- Homoscedacity is still intact.
- Normality is worse with the addition of this variable but it is still acceptable.
- No multicollinearity.

Model with log-transformed independent variables to see if the normality will be improved:

```
[150]: variables = variables = ['log_sqft_living', 'log_sqft_lot', 'condition', 'grade', 'has_view', 'waterfront', 'has_basement', 'floors', 'april', 'march', 'may', 'age<30', 'east', 'fareast', 'north', 'west']

y = data['log_price']
X = data[variables]

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
model.summary()
```

```
[150]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                OLS Regression Results
=====
Dep. Variable:                  log_price    R-squared:                  0.756
Model:                            OLS        Adj. R-squared:              0.756
Method:                 Least Squares    F-statistic:                 4044.
Date:                  Tue, 23 Aug 2022    Prob (F-statistic):          0.00
Time:                  20:19:03          Log-Likelihood:             -927.26
No. Observations:          20904          AIC:                        1889.
Df Residuals:              20887          BIC:                        2024.
Df Model:                   16
Covariance Type:           nonrobust
=====
===
                                coef      std err          t      P>|t|      [0.025

```

0.975]

-----					
---					
const	7.4255	0.044	170.017	0.000	7.340
7.511					
log_sqft_living	0.4525	0.007	60.360	0.000	0.438
0.467					
log_sqft_lot	0.0276	0.003	8.706	0.000	0.021
0.034					
condition	0.0796	0.003	26.573	0.000	0.074
0.086					
grade	0.1650	0.002	66.444	0.000	0.160
0.170					
has_view	0.1568	0.007	24.087	0.000	0.144
0.170					
waterfront	0.5308	0.024	22.115	0.000	0.484
0.578					
has_basement	-0.0048	0.005	-1.051	0.293	-0.014
0.004					
floors	0.0209	0.005	4.096	0.000	0.011
0.031					
april	0.0656	0.006	11.284	0.000	0.054
0.077					
march	0.0547	0.006	8.698	0.000	0.042
0.067					
may	0.0175	0.006	3.099	0.002	0.006
0.029					
age<30	-0.0201	0.005	-3.663	0.000	-0.031
-0.009					
east	0.4955	0.005	92.825	0.000	0.485
0.506					
fareast	0.4097	0.006	64.674	0.000	0.397
0.422					
north	0.3339	0.008	41.081	0.000	0.318
0.350					
west	0.5638	0.005	105.623	0.000	0.553
0.574					
=====					
Omnibus:	564.380	Durbin-Watson:	1.986		
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1424.854		
Skew:	0.003	Prob(JB):	3.95e-310		
Kurtosis:	4.279	Cond. No.	366.		
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
"""
```

```
[151]: variables = variables = ['log_sqft_living', 'log_sqft_lot', 'condition', 'grade', 'has_view', 'waterfront', 'floors', 'april', 'march', 'may', 'age<30', 'east', 'fareast', 'north', 'west']

y = data['log_price']
X = data[variables]

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
model.summary()
```

```
[151]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                        OLS Regression Results
=====
Dep. Variable:          log_price      R-squared:                0.756
Model:                  OLS           Adj. R-squared:            0.756
Method:                 Least Squares   F-statistic:              4313.
Date:                  Tue, 23 Aug 2022   Prob (F-statistic):       0.00
Time:                  20:19:03          Log-Likelihood:          -927.82
No. Observations:      20904            AIC:                     1888.
Df Residuals:          20888            BIC:                     2015.
Df Model:               15
Covariance Type:       nonrobust
=====
===

```

	coef	std err	t	P> t	[0.025	0.975]
const	7.4384	0.042	177.491	0.000	7.356	7.521
log_sqft_living	0.4489	0.007	67.003	0.000	0.436	0.462
log_sqft_lot	0.0285	0.003	9.297	0.000	0.022	0.034
condition	0.0796	0.003	26.573	0.000	0.074	0.086
grade	0.1652	0.002	66.689	0.000	0.160	0.170
has_view	0.1562	0.006	24.089	0.000	0.143	0.169

```

---

```

waterfront	0.5304	0.024	22.099	0.000	0.483
0.577					
floors	0.0230	0.005	4.859	0.000	0.014
0.032					
april	0.0656	0.006	11.284	0.000	0.054
0.077					
march	0.0547	0.006	8.700	0.000	0.042
0.067					
may	0.0175	0.006	3.096	0.002	0.006
0.029					
age<30	-0.0197	0.005	-3.602	0.000	-0.030
-0.009					
east	0.4955	0.005	92.819	0.000	0.485
0.506					
fareast	0.4100	0.006	64.788	0.000	0.398
0.422					
north	0.3335	0.008	41.080	0.000	0.318
0.349					
west	0.5627	0.005	107.414	0.000	0.552
0.573					


```
=====
Omnibus:                    564.179    Durbin-Watson:                1.986
Prob(Omnibus):              0.000    Jarque-Bera (JB):             1423.955
Skew:                      0.005    Prob(JB):                     6.20e-310
Kurtosis:                  4.279    Cond. No.                     351.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

""""

```
[152]: normality_homoscedecity(model);
```



- Slight improvement in R2 but slight worsening in Kurtosis and Jarque-Bera (JB).

- For ease of interpretation let's keep the previous model.

```
[153]: variables = ['sqft_living', 'sqft_lot', 'condition', 'grade', 'has_view',
    ↪, 'waterfront',
    'has_basement', 'floors',
    'april', 'march', 'may',
    'age<30',
    'east', 'fareast', 'north', 'west']

y = data['log_price']
X = data[variables]

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
model.summary()
```

```
[153]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  log_price    R-squared:                  0.753
Model:                            OLS       Adj. R-squared:              0.752
Method:                 Least Squares    F-statistic:                 3972.
Date:                Tue, 23 Aug 2022    Prob (F-statistic):          0.00
Time:                  20:19:03          Log-Likelihood:             -1067.8
No. Observations:          20904         AIC:                       2170.
Df Residuals:              20887         BIC:                       2305.
Df Model:                   16
Covariance Type:            nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	10.6384	0.019	563.259	0.000	10.601	10.675
sqft_living	0.0002	3.43e-06	62.667	0.000	0.000	0.000
sqft_lot	1.48e-06	1.57e-07	9.407	0.000	1.17e-06	1.79e-06
condition	0.0834	0.003	27.712	0.000	0.078	0.089
grade	0.1612	0.003	62.788	0.000	0.156	0.166
has_view	0.1433	0.007	21.805	0.000	0.130	0.156
waterfront	0.5184	0.024	21.474	0.000	0.471	0.566
has_basement	0.0171	0.004	4.005	0.000	0.009	0.026
floors	0.0300	0.005	6.257	0.000	0.021	0.039
april	0.0664	0.006	11.346	0.000	0.055	0.078
march	0.0547	0.006	8.651	0.000	0.042	0.067
may	0.0156	0.006	2.742	0.006	0.004	0.027
age<30	-0.0240	0.005	-4.440	0.000	-0.035	-0.013
east	0.4868	0.005	90.527	0.000	0.476	0.497
fareast	0.4012	0.006	62.808	0.000	0.389	0.414



north	0.3324	0.008	40.612	0.000	0.316	0.348
west	0.5316	0.005	105.152	0.000	0.522	0.541

```
=====
Omnibus:                    579.603    Durbin-Watson:                1.988
Prob(Omnibus):              0.000    Jarque-Bera (JB):            1406.417
Skew:                      -0.088    Prob(JB):                   3.99e-306
Kurtosis:                  4.258    Cond. No.                   2.24e+05
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.24e+05. This might indicate that there are strong multicollinearity or other numerical problems.

"""

```
[159]: coefs = model.params.apply('{0:.6f}'.format)
coefs = pd.DataFrame(coefs).reset_index()
coefs = coefs.rename({'index': 'var', 0: 'coef'}, axis = 1)
coefs.coef = pd.to_numeric(coefs.coef)
coefs
```

```
[159]:
```

	var	coef
0	const	10.638372
1	sqft_living	0.000215
2	sqft_lot	0.000001
3	condition	0.083434
4	grade	0.161232
5	has_view	0.143341
6	waterfront	0.518445
7	has_basement	0.017130
8	floors	0.029957
9	april	0.066388
10	march	0.054738
11	may	0.015626
12	age<30	-0.024027
13	east	0.486775
14	fareast	0.401208
15	north	0.332381
16	west	0.531562

```
[155]: # https://data.library.virginia.edu/
↳ interpreting-log-transformations-in-a-linear-model/
# Because the target variable is log-transformed let's exponentiate the_
↳ coefficients.
coefs['exp_coef'] = ((np.exp(coefs.coef) - 1) * 100).apply('{0:.6f}'.format)
coefs
```

```
[155]:
```

	var	coef	exp_coef
0	const	10.638372	4170382.047389
1	sqft_living	0.000215	0.021502
2	sqft_lot	0.000001	0.000100
3	condition	0.083434	8.701347
4	grade	0.161232	17.495753
5	has_view	0.143341	15.412329
6	waterfront	0.518445	67.941413
7	has_basement	0.017130	1.727756
8	floors	0.029957	3.041023
9	april	0.066388	6.864127
10	march	0.054738	5.626384
11	may	0.015626	1.574872
12	age<30	-0.024027	-2.374065
13	east	0.486775	62.706048
14	fareast	0.401208	49.362791
15	north	0.332381	39.428397
16	west	0.531562	70.158811

#### Model summary:

- R-squared of 0.753 means that the dependent variables explain 75% of the variability in price.
- All variables are statistically significant ( $p < .05$ ) meaning we can reject the null that are not related to price.
- Normality is acceptable, homoscedasticity is preserved, no presence of multicollinearity,
- Durbin-Watson score is between 1.5 and 2.5, meaning: no first-order autocorrelation.
- Skewness is between -0.5 and 0.5, it is approximately symmetric.
- Kurtosis is 4.2 which is not ideal but acceptable. An increased kurtosis ( $>3$ ) can be visualized as a thinner “bell” with a higher peak. Hair et al. (2010) and Bryne (2010) argues that data is considered normal if skewness is between -2 to +2 and kurtosis is between -7 to +7.

#### Coefficient interpretation:

- For every 1 unit increase in `sqft_living` price increases by about 0.022 %
- For every 100 sqft increase in the house price increases by about 2.2 %.
- Average sqft of a house is 2000. Given all other variables are kept constant, if you increase a house size from 2000 to 3000 sqft you would increase price by 22%.
- For every 1 unit increase in `sqft_lot` price increases by 0.0001 %.
- For every 10000 sqft increase in the lot, price increases by about 1 %.
- For every 1 unit increase in `grade` price increases by 17 %.
- Being `waterfront` increases price by 68% compared to being non waterfront (given all other factors are same).
- Having `view` increases price by 15% compared to not having a view (given all other factors are same).

- For every `floor` added, price increases by 3%.
- Houses sold in `April` are 6.8% more expensive than those sold in winter-fall or summer.
- Houses sold in `March` are 5.6% more expensive than those sold in winter-fall or summer.
- Houses sold in `May` are 1.5% more expensive than those sold in winter-fall or summer.
- Houses in `West-Seattle` are 70% more expensive than those in South (given all other factors are the same).
- Houses in `East` area are 62% more expensive than those in South.

#### 0.11.9 Recommendations based on regression results:

1. Invest on increasing the total square footage of the house as much as possible (rather than investing on the lot size). For every 1000 sqft increase in the house price increases by about 22%.
2. Being on waterfront increases the house price by 67%, so invest on houses on waterfront. Having a view increases the house price by 15%, so invest on houses that has a view.
3. Put the house on the market in April which increases the price by 6.8%. The next best month is March with a 5.6% increase.
4. Invest on houses in Seattle for 70% increase in price, and Medina, Bellevue, Mercer Island and Kirkland for a 62% increase compared to the South.

#### 0.11.10 Limitations

- Skewed data required outlier removal or data transformation which made the interpretations trickier.
- Although Zipcodes provided very useful information with regard to price, I did not feel comfortable using all 70 categories in the regression model.
- City to Zipcode mapping did not work due to the zipcode and city boundaries not being the same - many zipcodes belonging to various cities and vice versa. Total number of cities were also more than desired: 32.

#### 0.11.11 Improvements

- Gather more detailed location info using API calls.
- Using cluster analysis group zipcodes into more meaningful clusters.

```
[160]: ./nbconvert.py --format=pdf KingCountySales.ipynb
```

```
File "<ipython-input-160-745bf5d3884a>", line 1
  ./nbconvert.py --format=pdf KingCountySales.ipynb
  ~
SyntaxError: invalid syntax
```

```
[ ]:
```